

OPERATING SYSTEMS: FILE SYSTEMS



Files, directories and file system

To remember...

Before classes

Class

After class

Prepare the prerequisites.

Study the material associated with the **bibliography**:
slides alone are not enough.
Please ask questions (especially after study).

Exercising skills:

- ▶ Perform all **exercises**.
- ▶ Carrying out the **practice notebooks** and **the practical exercises** progressively.

Recommended reading

Base



1. Carretero 2020:
 1. Cap. 6
2. Carretero 2007:
 1. Cap. 9.1-9.5,
 2. Cap. 9.8-9.10 & 9.12

Suggested



1. Tanenbaum 2006:
 1. (es) Cap. 6
 2. (en) Cap. 6
2. Stallings 2005:
 1. 12.1-12.8
3. Silberschatz 2006:
 1. 10.3-10.4,
 2. 11.1-11.6 and 13

Contents

- Introduction
- File
- Directory
- File System
- Partitions/Volumes
- Devices
- System software
- File System (manager)

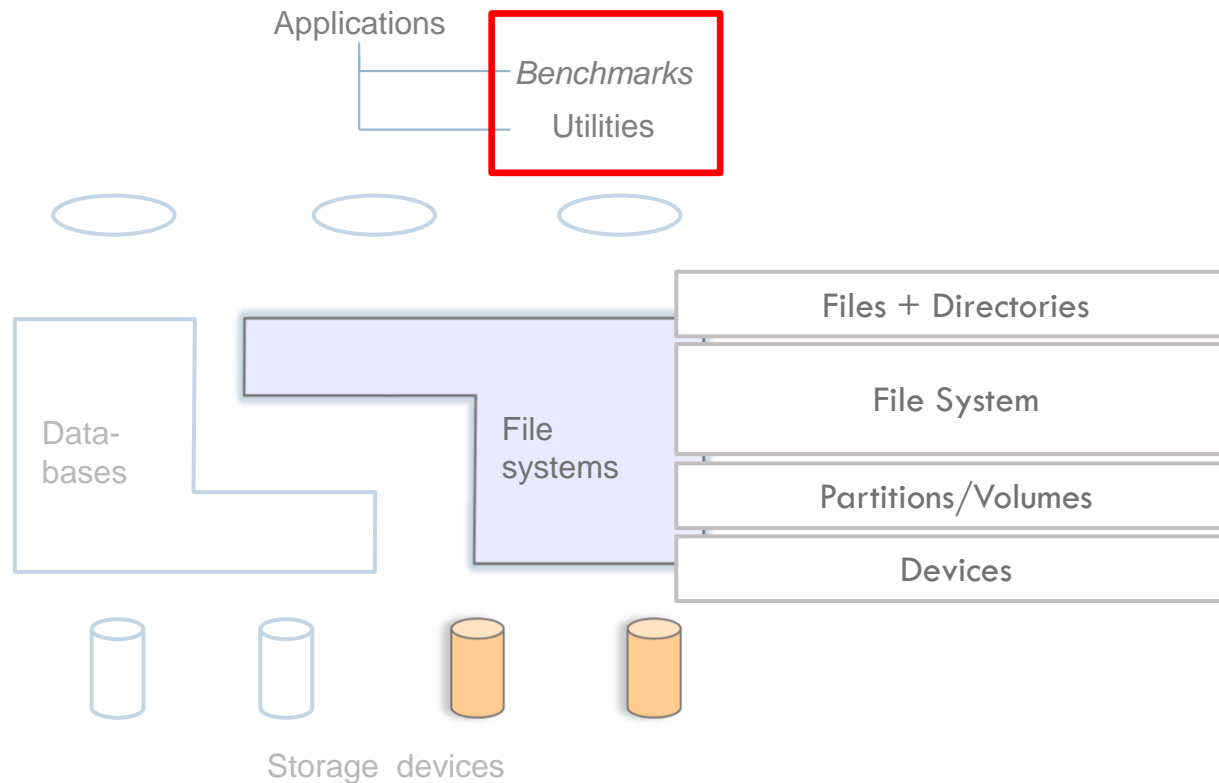
Contents

- Introduction
- File
- Directory
- File System
- Partitions/Volumes
- Devices
- **System software**
- File System (manager)

System software

6

Alejandro Calderón Mateos 



Benchmarks

□ Benchmarks:

- They allow to measure the **performance of the file system** (and any dependency on it)
- Designed to measure **different aspects**:
latency, bandwidth,
number of files processed per unit time, etc.
- Examples working with metadata: fdtree, mdtest, etc.
- Examples working with data: iozone, postmark, IOR, etc.

File system consistency

- Software failures may result in inconsistent information (and metadata).
- Solution:
 - ▣ Availability of tools to check the file system and repair the errors found.
- Two important aspects to review:
 - ▣ Verify that the **physical structure** of the file system is coherent
 - ▣ Verify that the **logical structure** of the file system is correct.

File system consistency

- Software failures may result in inconsistent information (and metadata).
- Solution:
 - ▣ Availability of tools to check the file system and repair the errors found.
- Two important aspects to review:

- ▣ Verify that the **physical structure** of the file system is coherent

 - ▣ Verify that the **logical structure** of the file system is correct.

File system consistency

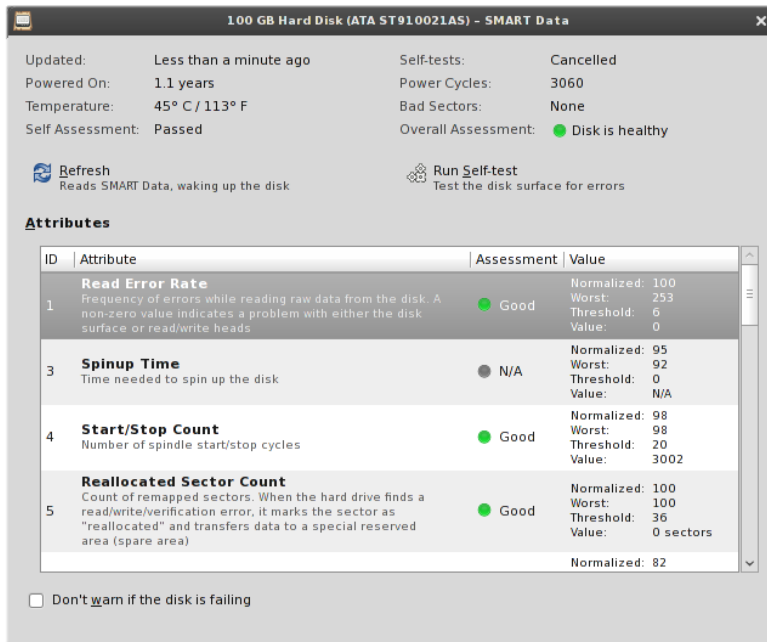
physical structure

10

Alejandro Calderón Mateos 



□ Controller logic:

- Disk-controller status tests are performed.
- E.g.: S.M.A.R.T.



100 GB Hard Disk (ATA ST910021AS) - SMART Data

Updated: Less than a minute ago Self-tests: Cancelled
Powered On: 1.1 years Power Cycles: 3060
Temperature: 45° C / 113° F Bad Sectors: None
Self Assessment: Passed Overall Assessment: ● Disk is healthy

 Refresh Reads SMART Data, waking up the disk  Run Self-test Test the disk surface for errors


Attributes

| ID | Attribute | Assessment | Value |
|--|-----------|---|--|
| Read Error Rate Frequency of errors while reading raw data from the disk. A non-zero value indicates a problem with either the disk surface or read/write heads | | | |
| 1 | | ● Good | Normalized: 100 Worst: 253 Threshold: 6 Value: 0 |
| Spinup Time Time needed to spin up the disk | | | |
| 3 | | ● N/A | Normalized: 95 Worst: 92 Threshold: 0 Value: N/A |
| Start/Stop Count Number of spindle start/stop cycles | | | |
| 4 | | ● Good | Normalized: 98 Worst: 98 Threshold: 20 Value: 3002 |
| Reallocated Sector Count Count of remapped sectors. When the hard drive finds a read/write/verification error, it marks the sector as "reallocated" and transfers data to a special reserved area (spare area) | | | |
| 5 | | ● Good | Normalized: 100 Worst: 100 Threshold: 36 Value: 0 sectors |
| | | | Normalized: 82 |

☐ Don't warn if the disk is failing

□ Disk surface:

- Reads/writes disk blocks one by one to check for problems on the surface of part of the disk.
- E.g.: if what is read is different from what is written



File system consistency

- Software failures may result in inconsistent information (and metadata).
- Solution:
 - ▣ Availability of tools to check the file system and repair the errors found.
- Two important aspects to review:
 - ▣ Verify that the **physical structure** of the file system is coherent
 - ▣ Verify that the **logical structure** of the file system is correct.

File system consistency

logical structure

12

Alejandro Calderón Mateos 

- Disk structures:
 - ▣ Check that the data structure on disk is consistent for partition, directories and files
 - ▣ E.g.: fsck in Linux, scandisk in Windows

```
acaldero@phoenix:/tmp$ sudo fsck -f /dev/loop1
fsck desde util-linux-ng 2.17.2
e2fsck 1.41.12 (17-May-2010)
Paso 1: Verificando nodos-i, bloques y tamaños
Paso 2: Verificando la estructura de directorios
Paso 3: Revisando la conectividad de directorios
Paso 4: Revisando las cuentas de referencia
Paso 5: Revisando el resumen de información de grupos
/dev/loop1: 11/28560 ficheros (0.0% no contiguos), 5161/114180 bloques
acaldero@phoenix:/tmp$
```

File system consistency

logical structure

13

Alejandro Calderón Mateos 

- File System on disk:
 - Check that the content of the superblock corresponds to the characteristics of the file system.
 - It is checked that the i-node bitmaps correspond to the occupied i-nodes in the file system.
 - Check that the bitmaps of blocks correspond to the blocks assigned to files.
 - Check that no block is assigned to more than one file.
- Directories:
 - The directory system of the file system is checked to see that the same node-i is not assigned to more than one directory.
- Files:
 - The protection and privilege bits are checked.
 - The link counter is checked.

Backup

Where?

14

Alejandro Calderón Mateos 

□ Place:

- Distant from the main system
- Protected from water, fire, etc.
 - Fireproof cabinets



□ Medium:

- Hard disk
 - A: capacity and price, D: fragile
- Tape
 - A: capacity and price, D: slow






Backup

How?

15

Alejandro Calderón Mateos 

- **Full backup:**
copy the entire contents of the file system.
- **Differential backup:**
contains all files that have been changed since the last **full backup**.
- **Incremental backup:**
contains all files that have been modified since the last **full backup** or **differential backup**



Backup

When?

16

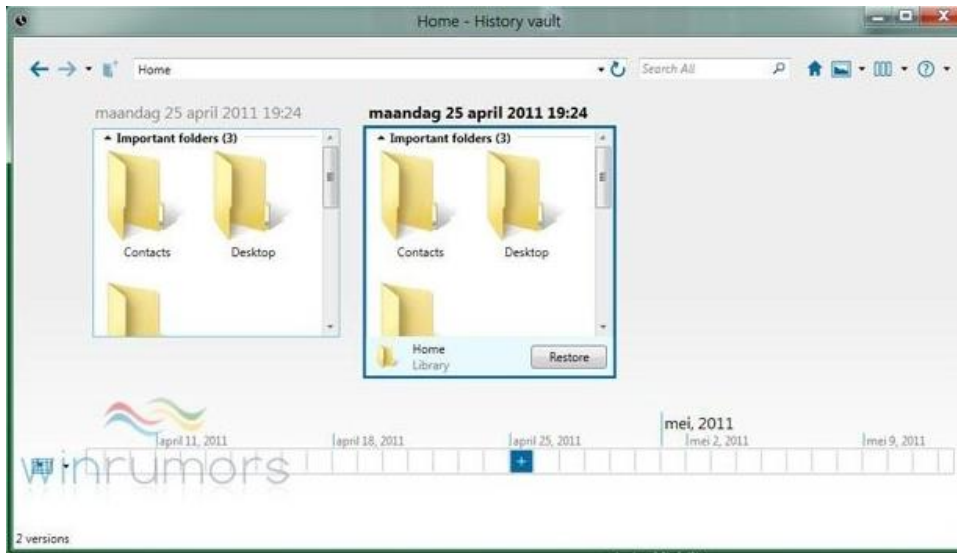
Alejandro Calderón Mateos 

- **Off-line:**
 - ▣ The backup is performed during periods of time when the system data is not in use.
- **On-line:**
 - ▣ The backup is performed while the system is in use.
 - ▣ Use of techniques to avoid consistency problems:
 - *Snapshots*
read-only copy of the file system state.
 - *Copy-on-write*
writes after snapshot are performed in copies.

Backup copy

17

Alejandro Calderón Mateos 



<http://www.genbeta.com/systems-operativos/primeras-imagenes-de-history-vault-el-time-machine-de-windows-8>



http://www.reghardware.com/2007/11/08/review_leopard_pt2/page2.html

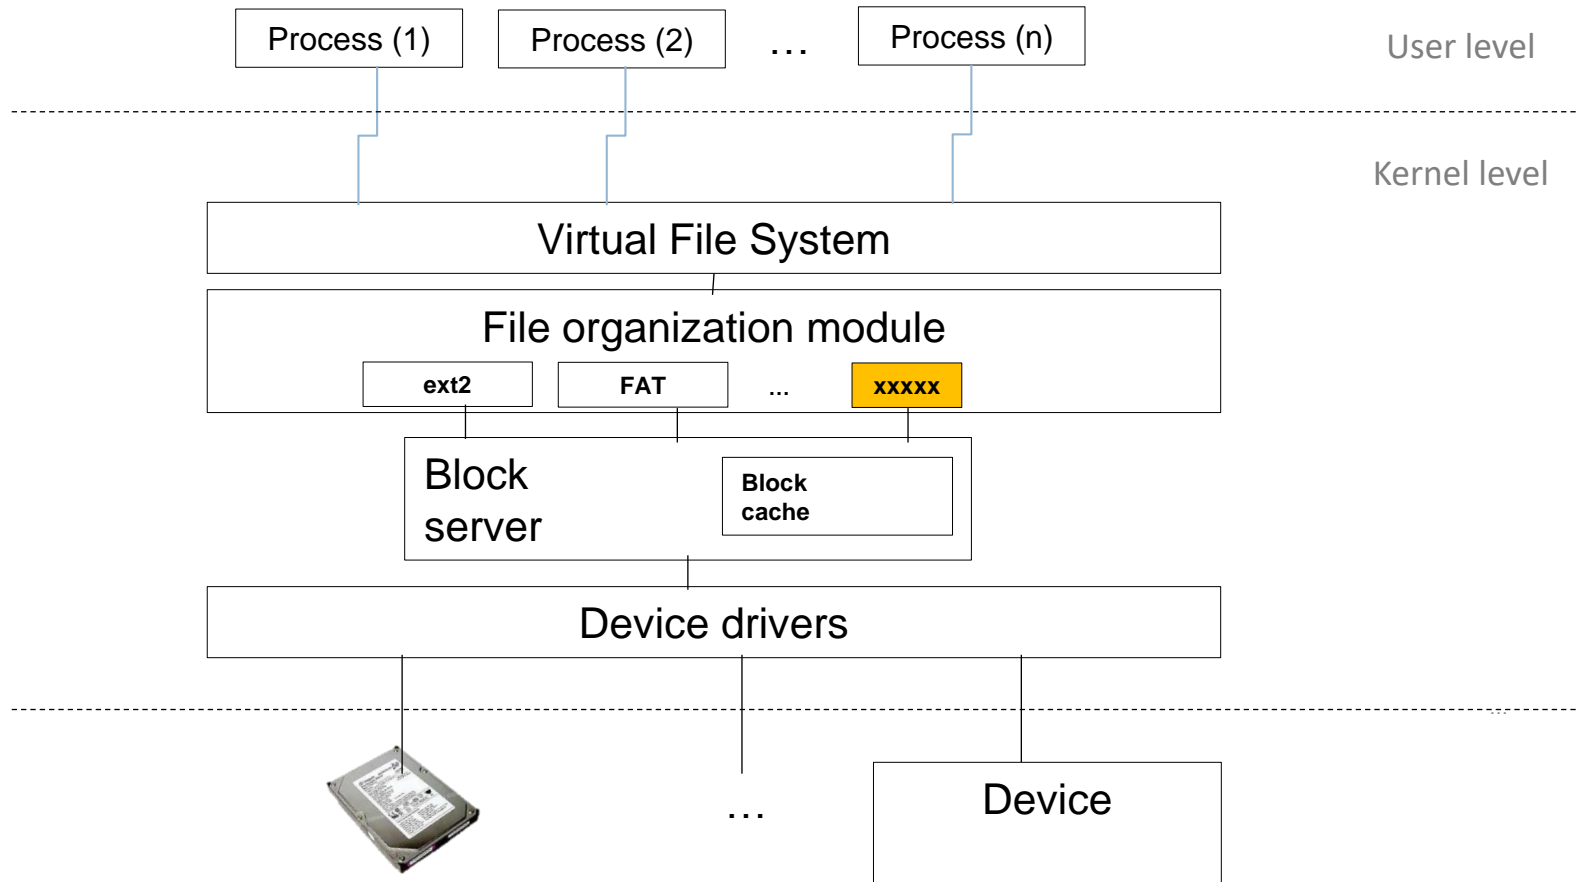
Contents

- Introduction
- File
- Directory
- File System
- Partitions/Volumes
- Devices
- System software
- **File System (manager)**

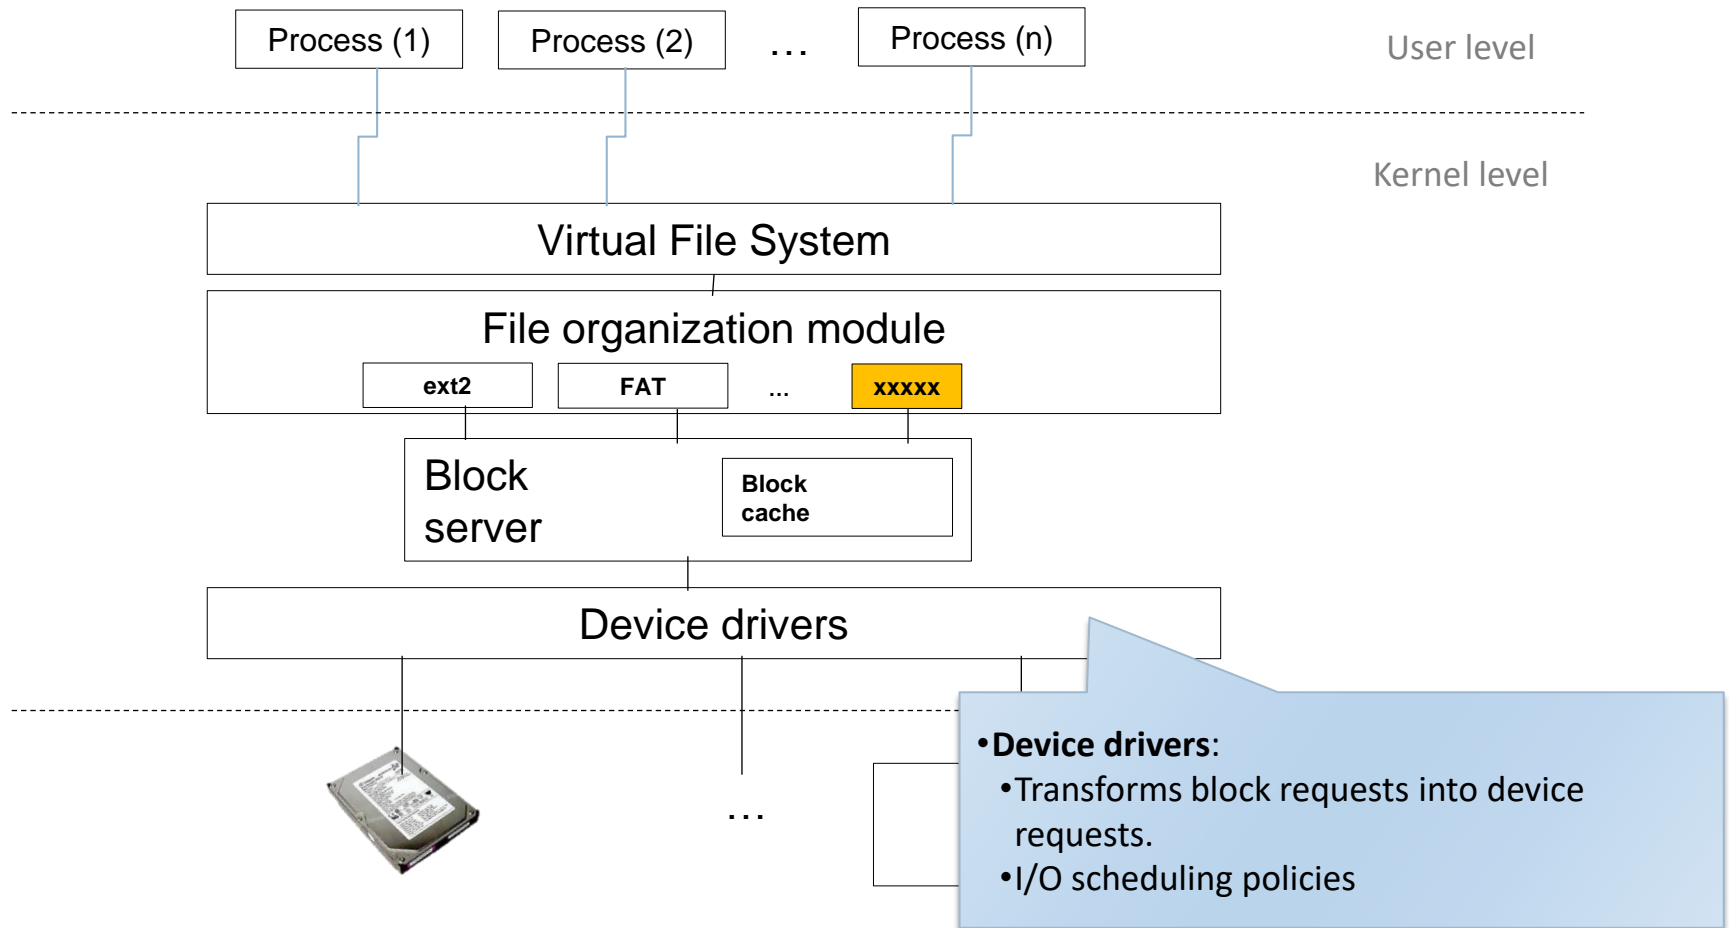
File management architecture...

19

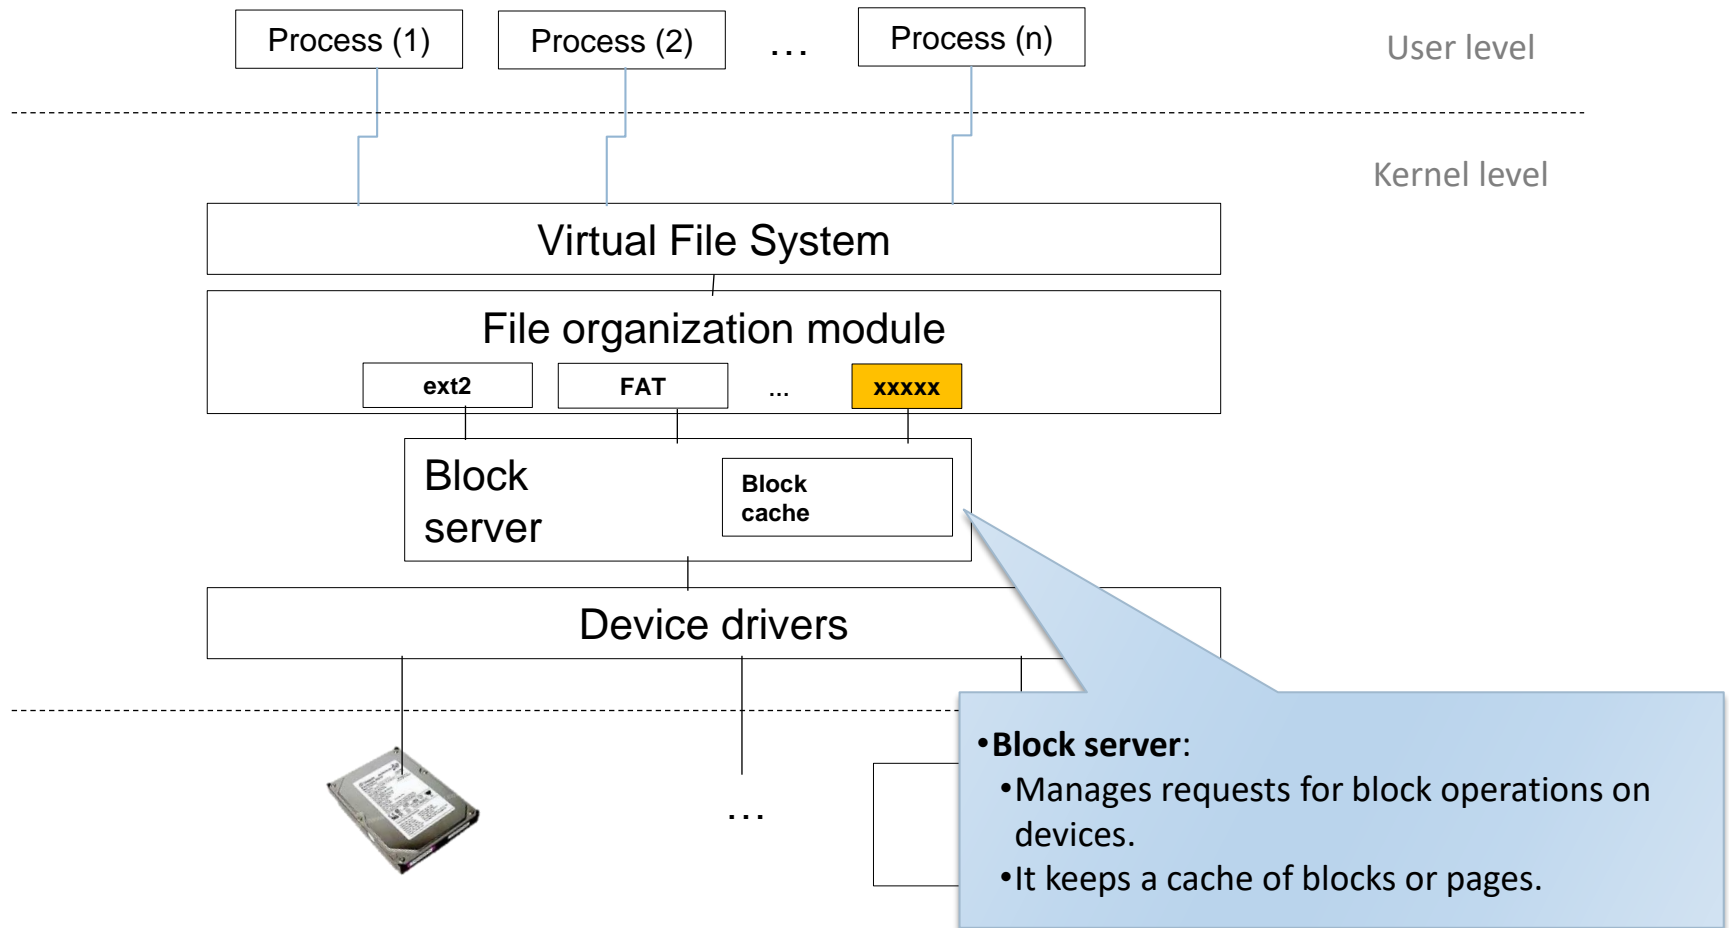
Alejandro Calderón Mateos 



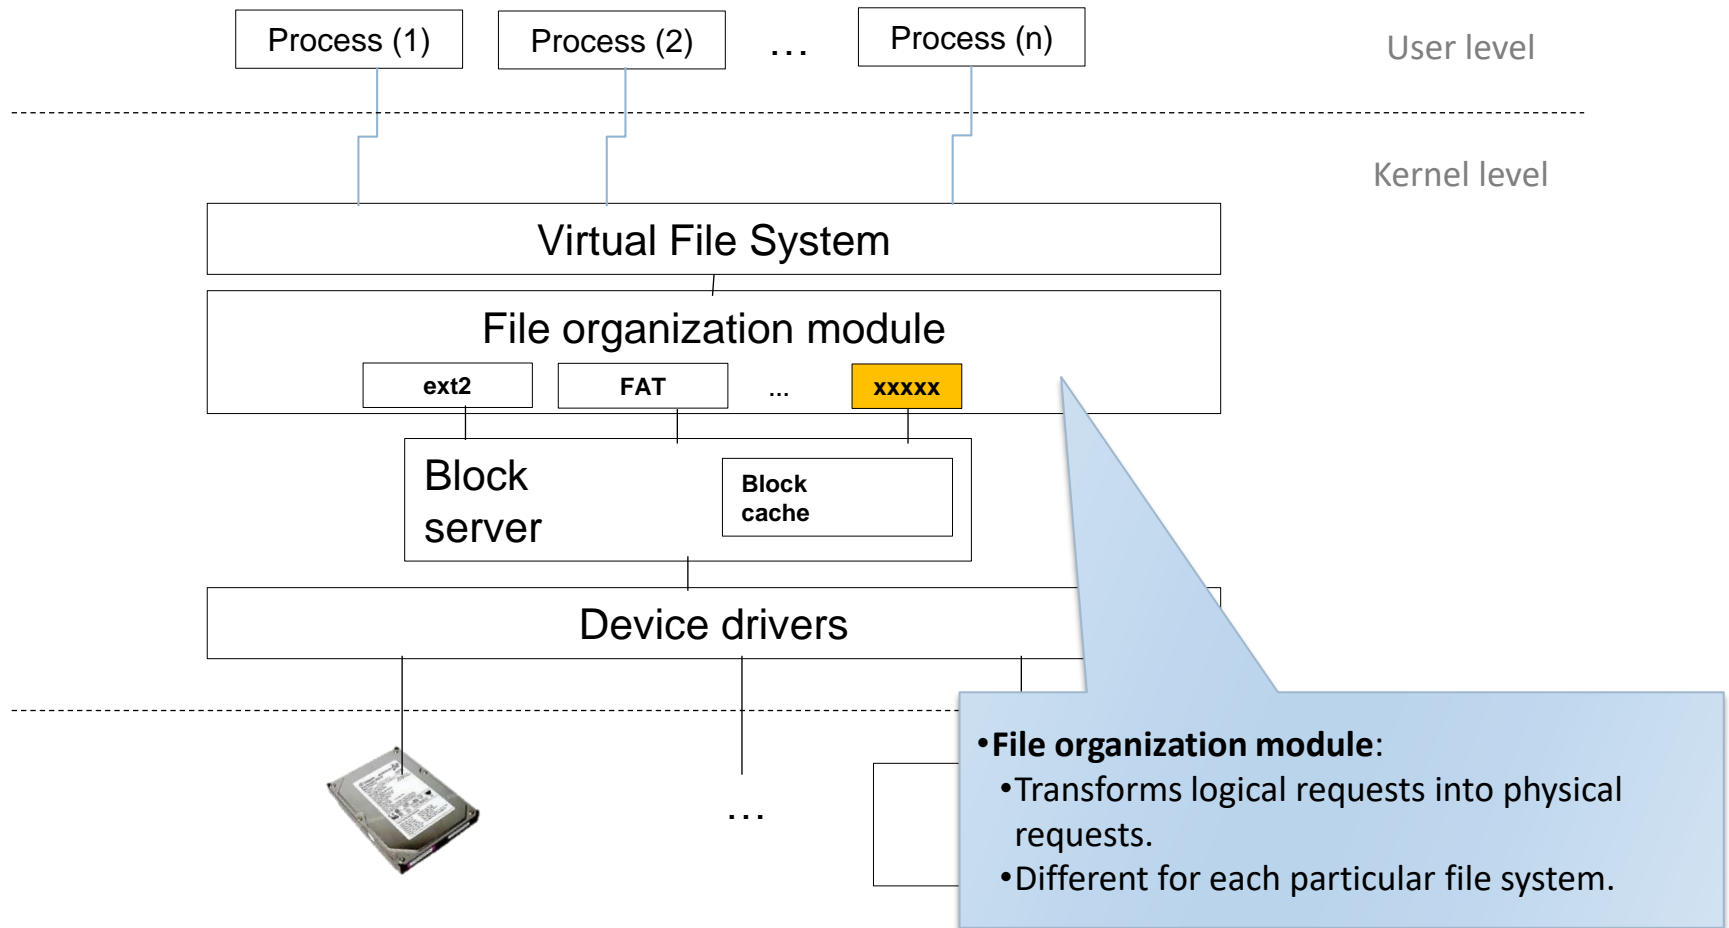
File management architecture...



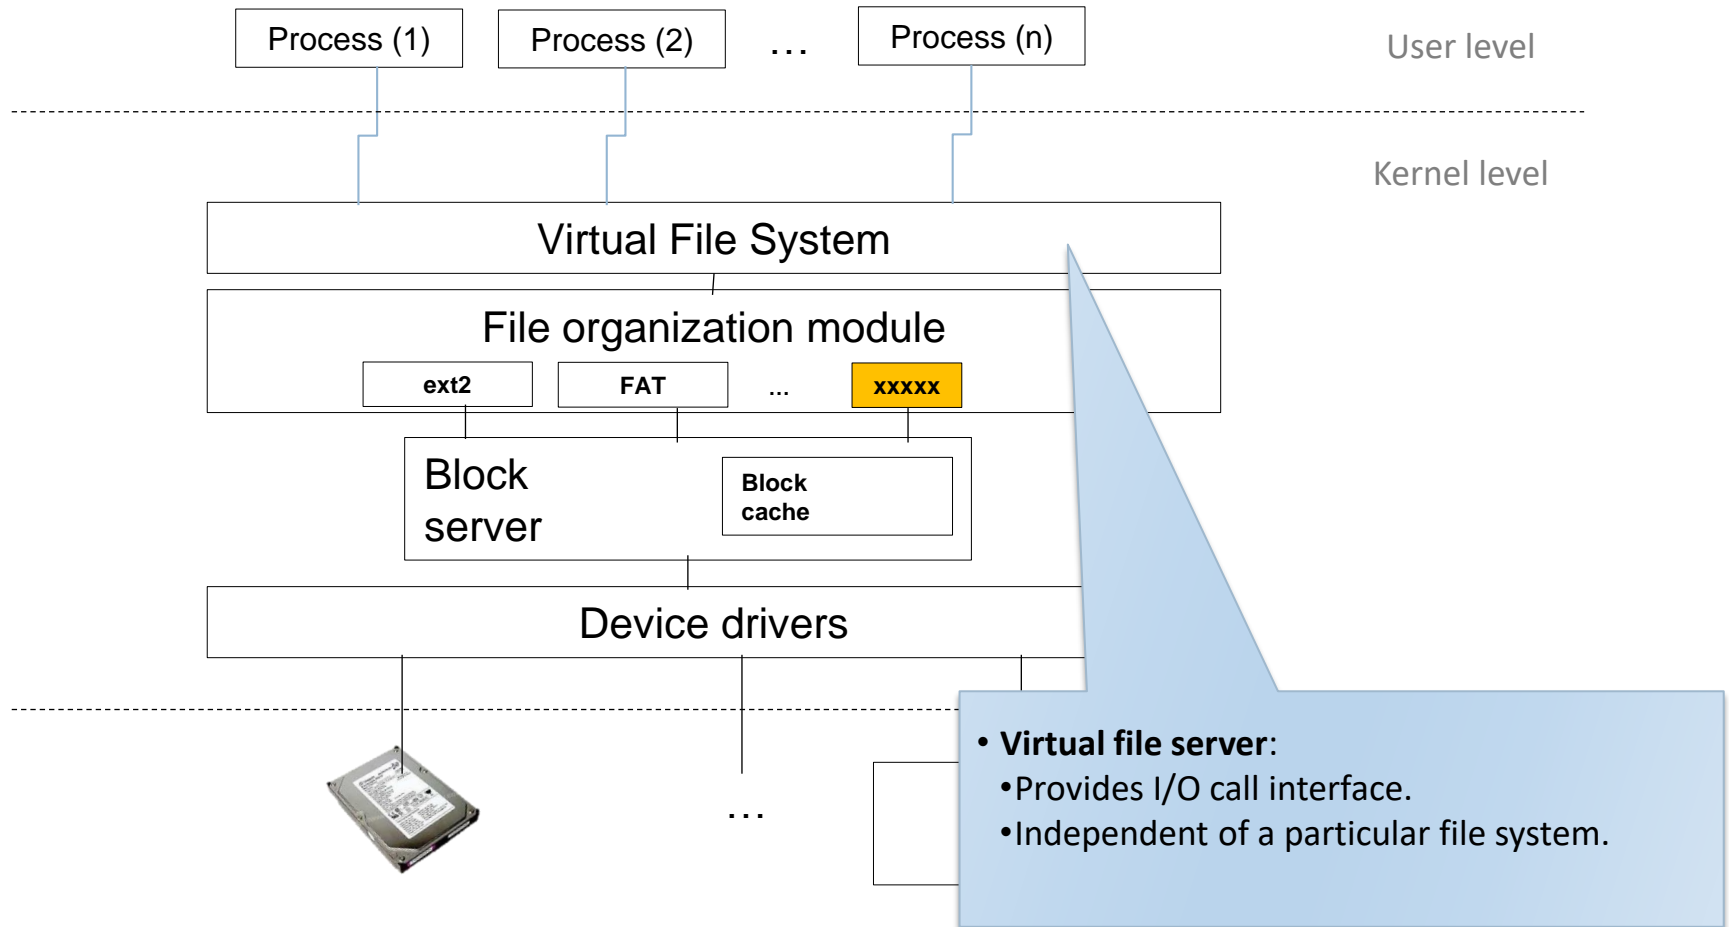
File management architecture...



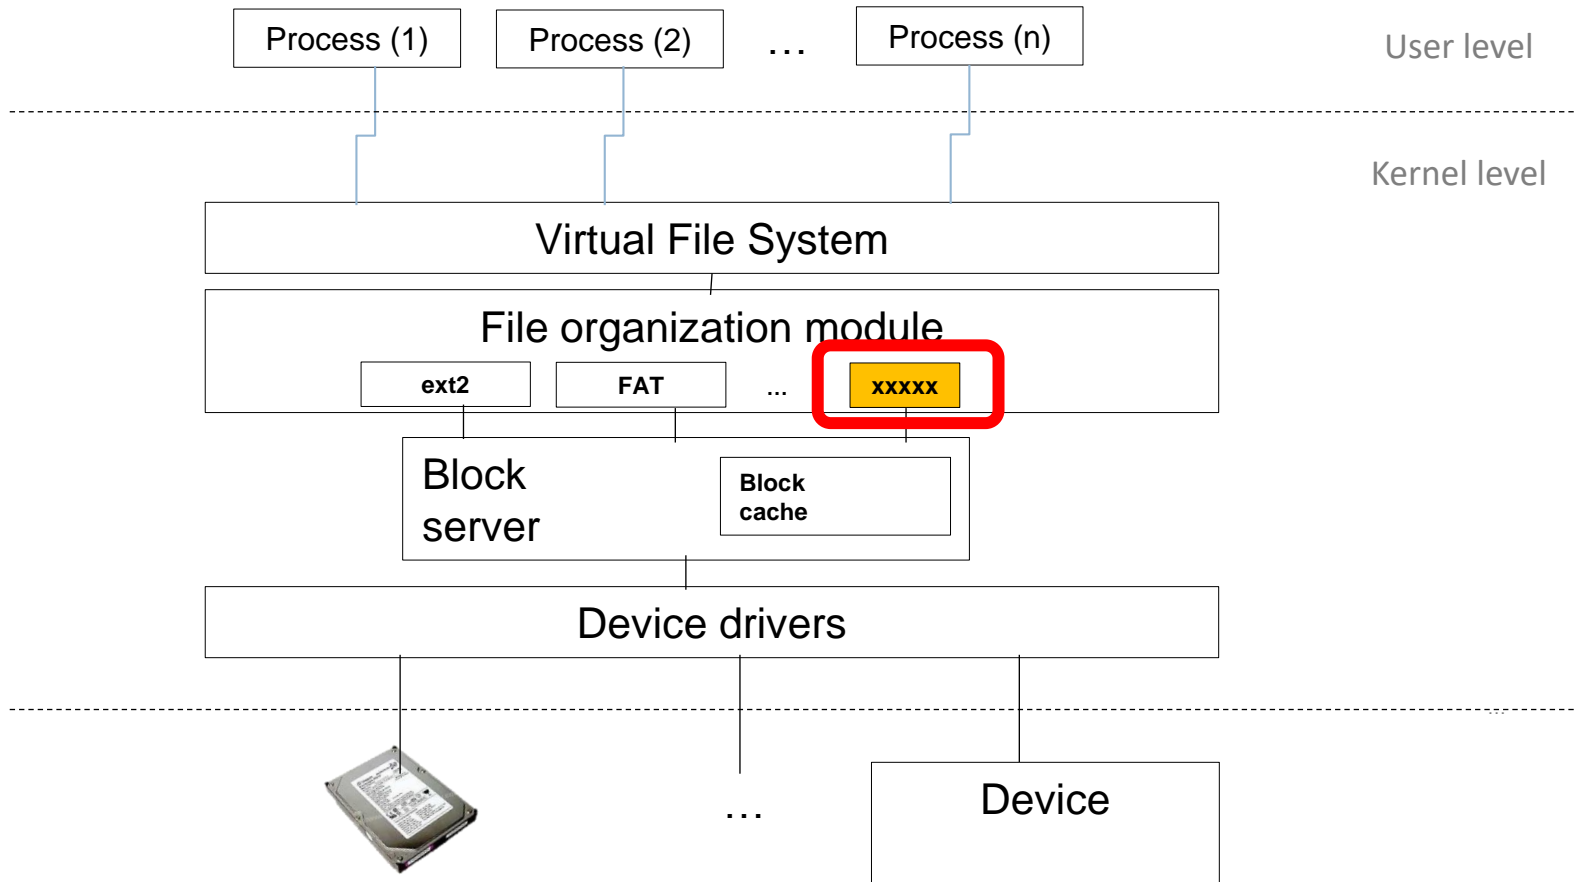
File management architecture...



File management architecture...



Destination (related to architecture)... file system design and implementation

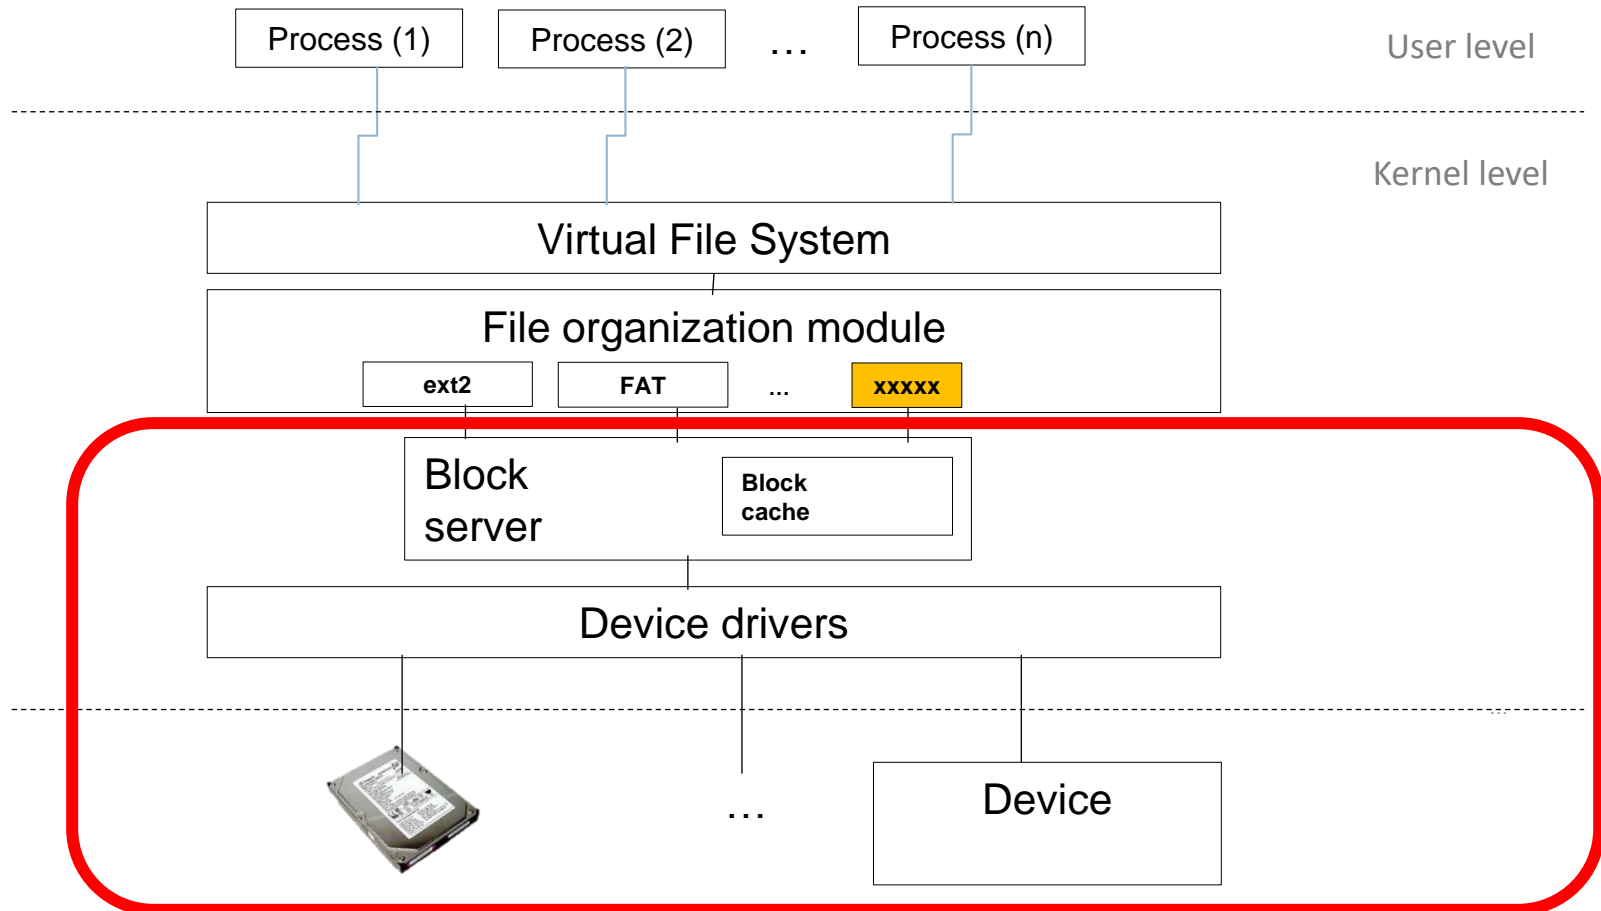


Origin (related to architecture)...

a) disk blocks + b) disk block cache

25

Alejandro Calderón Mateos 

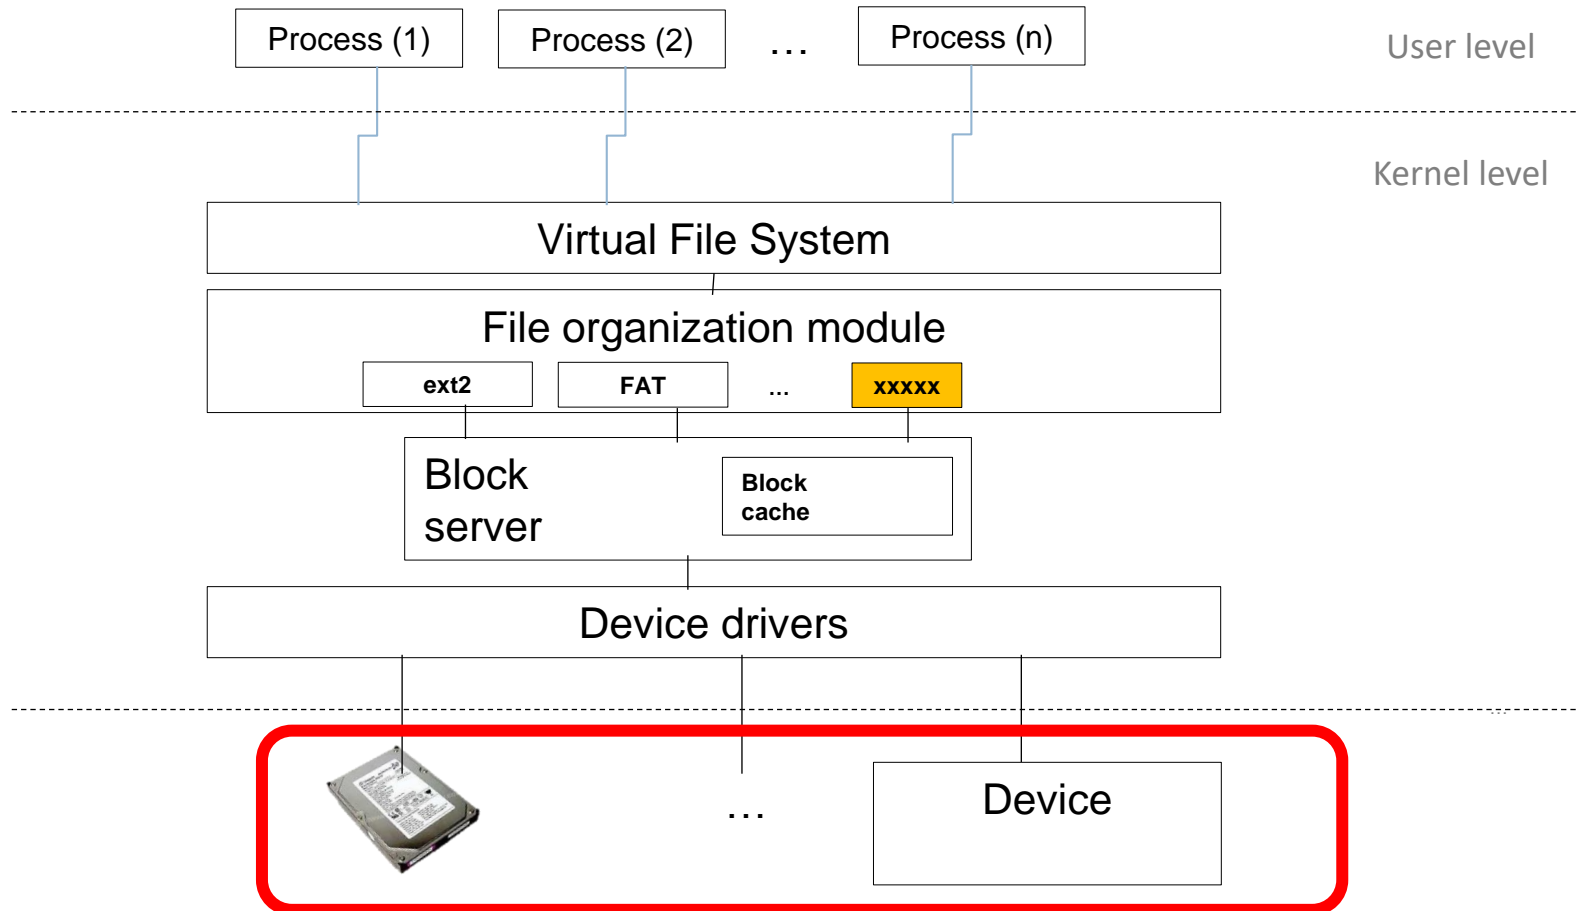


Origin (related to architecture)...

a) disk blocks

26

Alejandro Calderón Mateos 

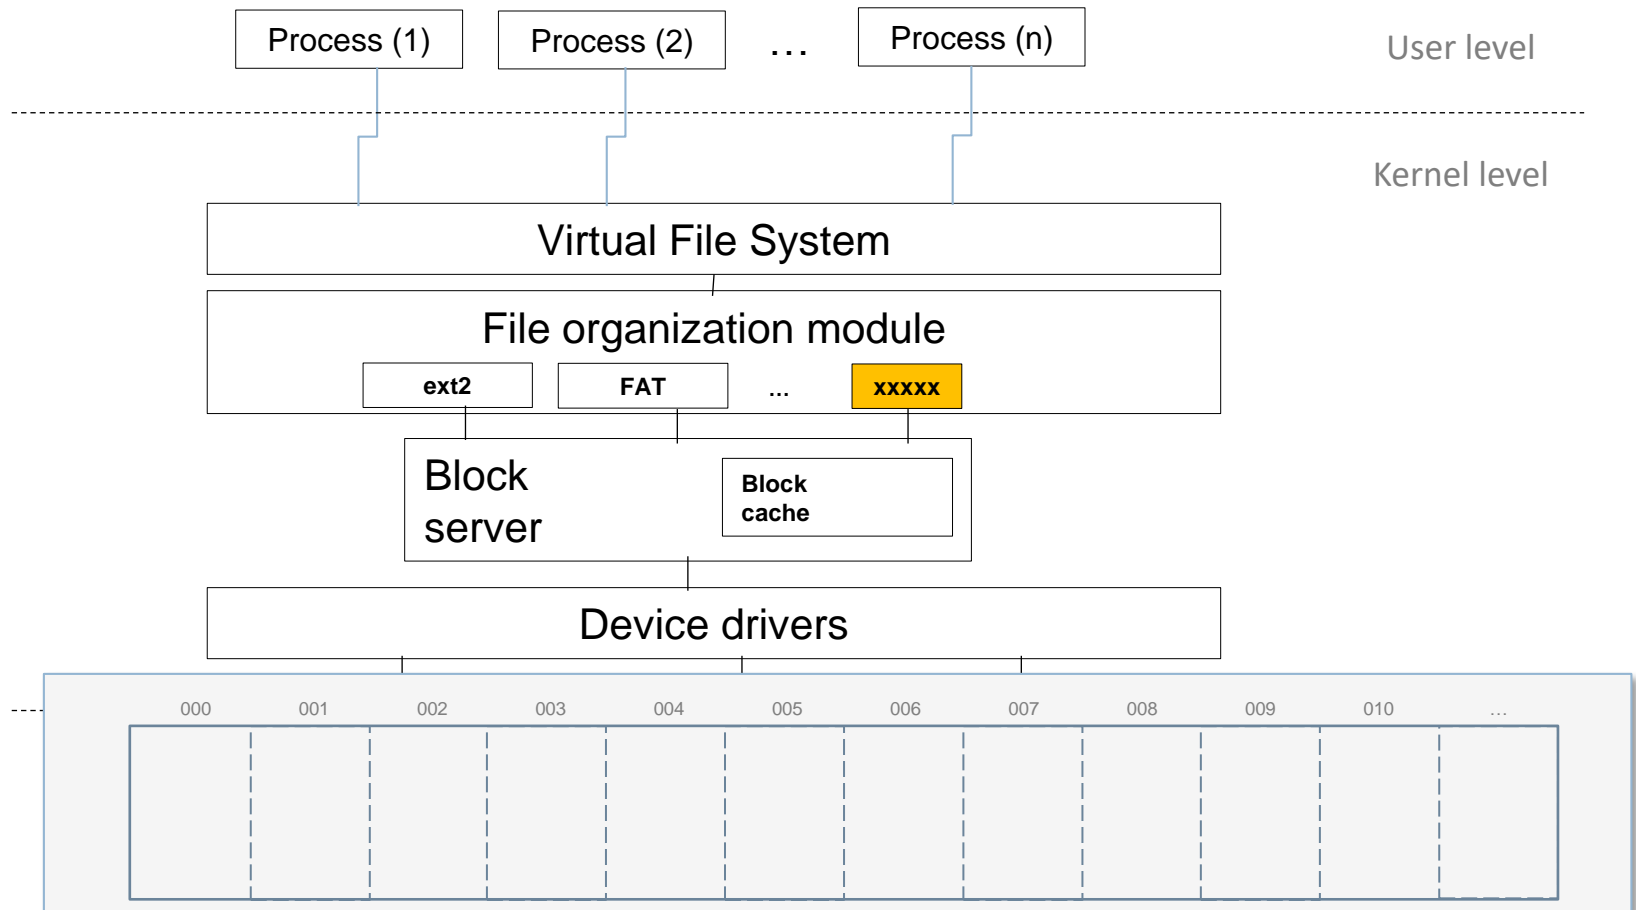


Origin (related to architecture)...

a) disk blocks

27

Alejandro Calderón Mateos 

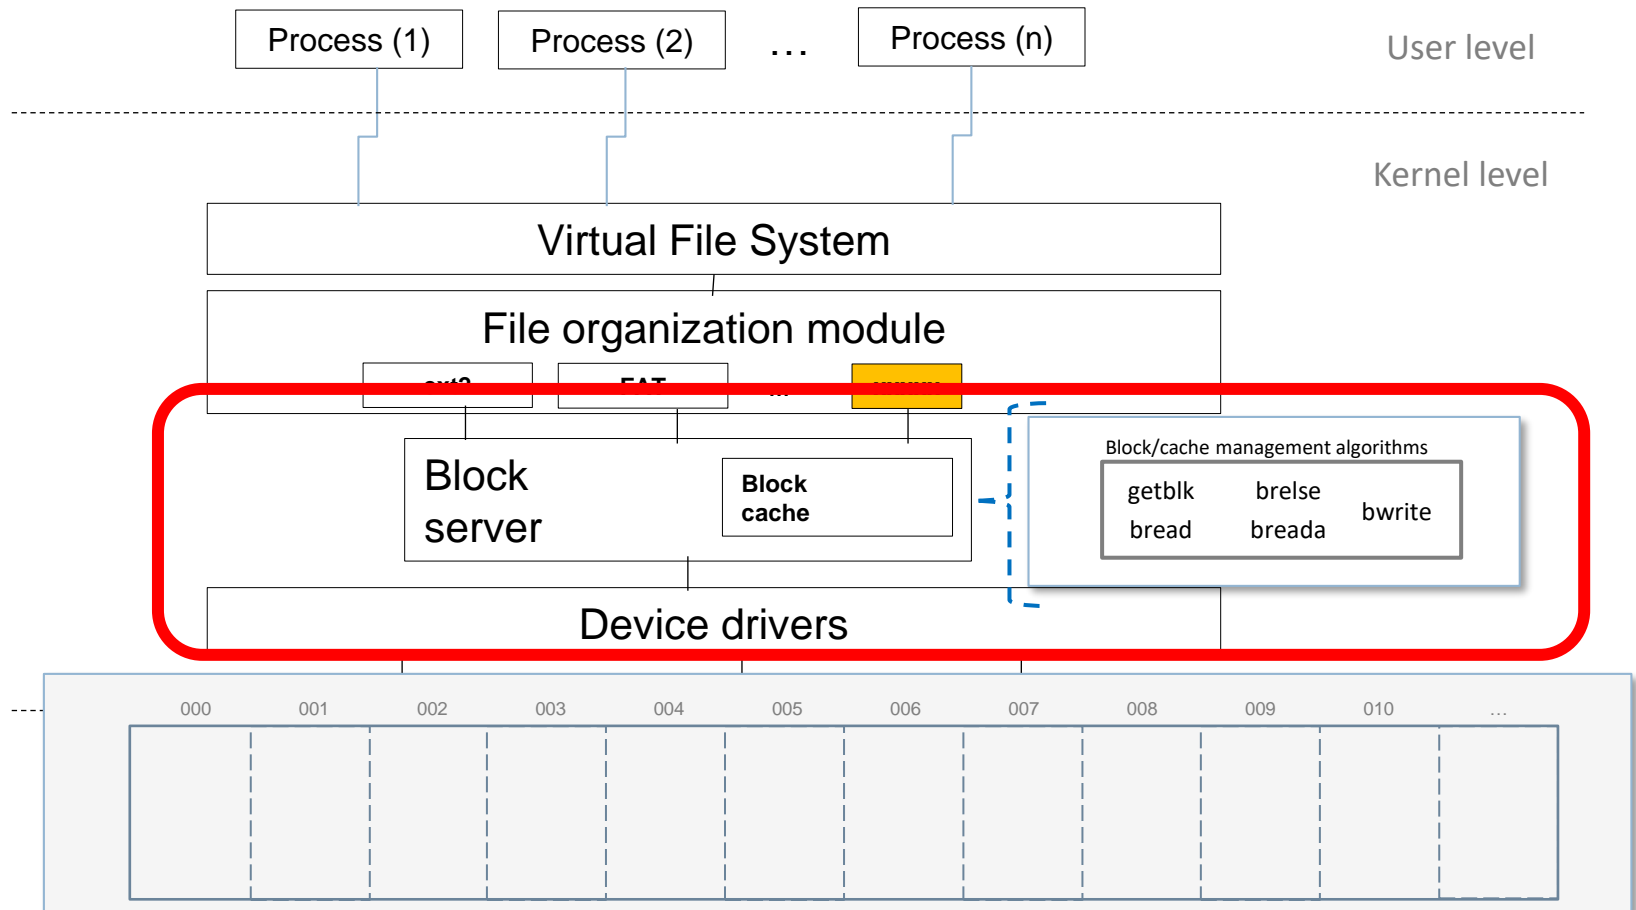


Origin (related to architecture)...

b) disk block cache

28

Alejandro Calderón Mateos



Origin (related to architecture)...

b) disk block cache

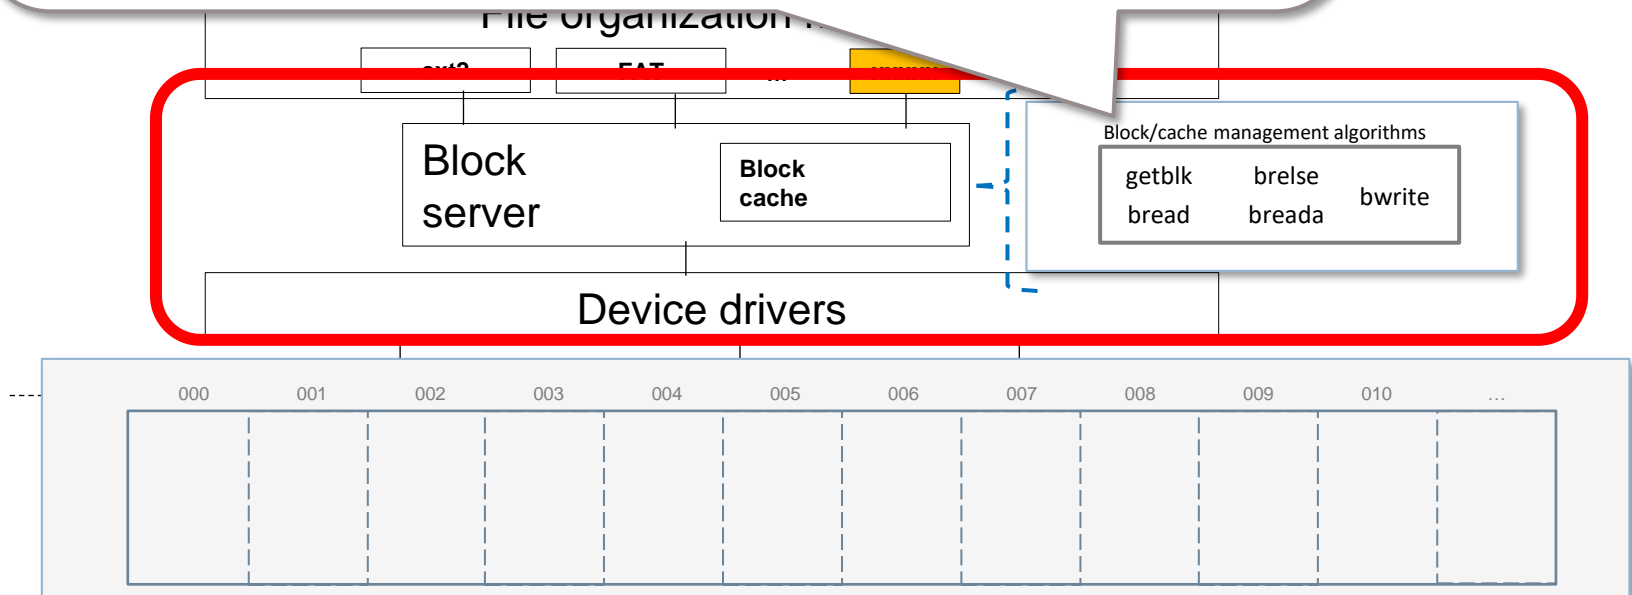
29

Alejandro Calderón Mateos

- ▶ **getblk**: searches/allocates a block in cache (from a given v-node, offset and size).
- ▶ **brelease**: releases a block and adds it to the free list.
- ▶ **bwrite**: writes a cache block to disk.
- ▶ **bread**: reads a block from disk to cache.
- ▶ **breada**: reads 1 block (and the next) from disk to cache.

User level

Kernel level



Block server

- It is in charge of:
 - ▣ Issue generic commands to read and write blocks to device handlers (using the device-specific routines).
 - ▣ Optimize I/O requests.
 - Ej.: block cache.
 - Can be integrated with virtual memory manager.
 - ▣ Provide a logical naming for the devices.
 - E.g.: /dev/hda3 (third partition of the first disk)

Block server

- General operation:
 - If the block is in cache
 - Copy content (+ update block usage metadata)
 - If the block is not in cache
 - Read the device block and store it in the cache
 - Copy content (and update metadata)
 - If the block has been written on (dirty)
 - Writing policy
 - If the cache is full, it is necessary to make room for it
 - Replacement policy

Block server

□ General operation:

○ **read-ahead:**

- Read a number of blocks after the required one and cached (improves performance on consecutive accesses)

- Read the device block and store it in the cache
- Copy content (and update metadata)
- If the block has been written on (dirty)
 - Writing policy
- If the cache is full, it is necessary to make room for it
 - Replacement policy

Block server

- **write-through:**
 - It is written each time the block is modified (– yield, + reliability)
- **write-back:**
 - Data are only written to disk when they are chosen for replacement due to lack of cache space (+ performance, – reliability)
- **delayed-write:**
 - Write to disk the modified data blocks in the cache periodically every certain time (30 seconds in UNIX) (compromise between previous)
- **write-on-close:**
 - When a file is closed, its blocks are dumped to disk..

■ If the data has been written on (dirty)

■ Writing policy

■ If the cache is full, it is necessary to make room for it

■ Replacement policy

Block server

□ General operation:

□ If the block is in cache

- Copy content (+ update block usage metadata)

□ If the block is not in cache

- Read the device block and store it in the cache

- **FIFO** (*First in First Out*)
- **Clock algorithm** (*Second opportunity*)
- **MRU** (*Most Recently Used*)
- **LRU** (*Least Recently Used*) <- + frequently used

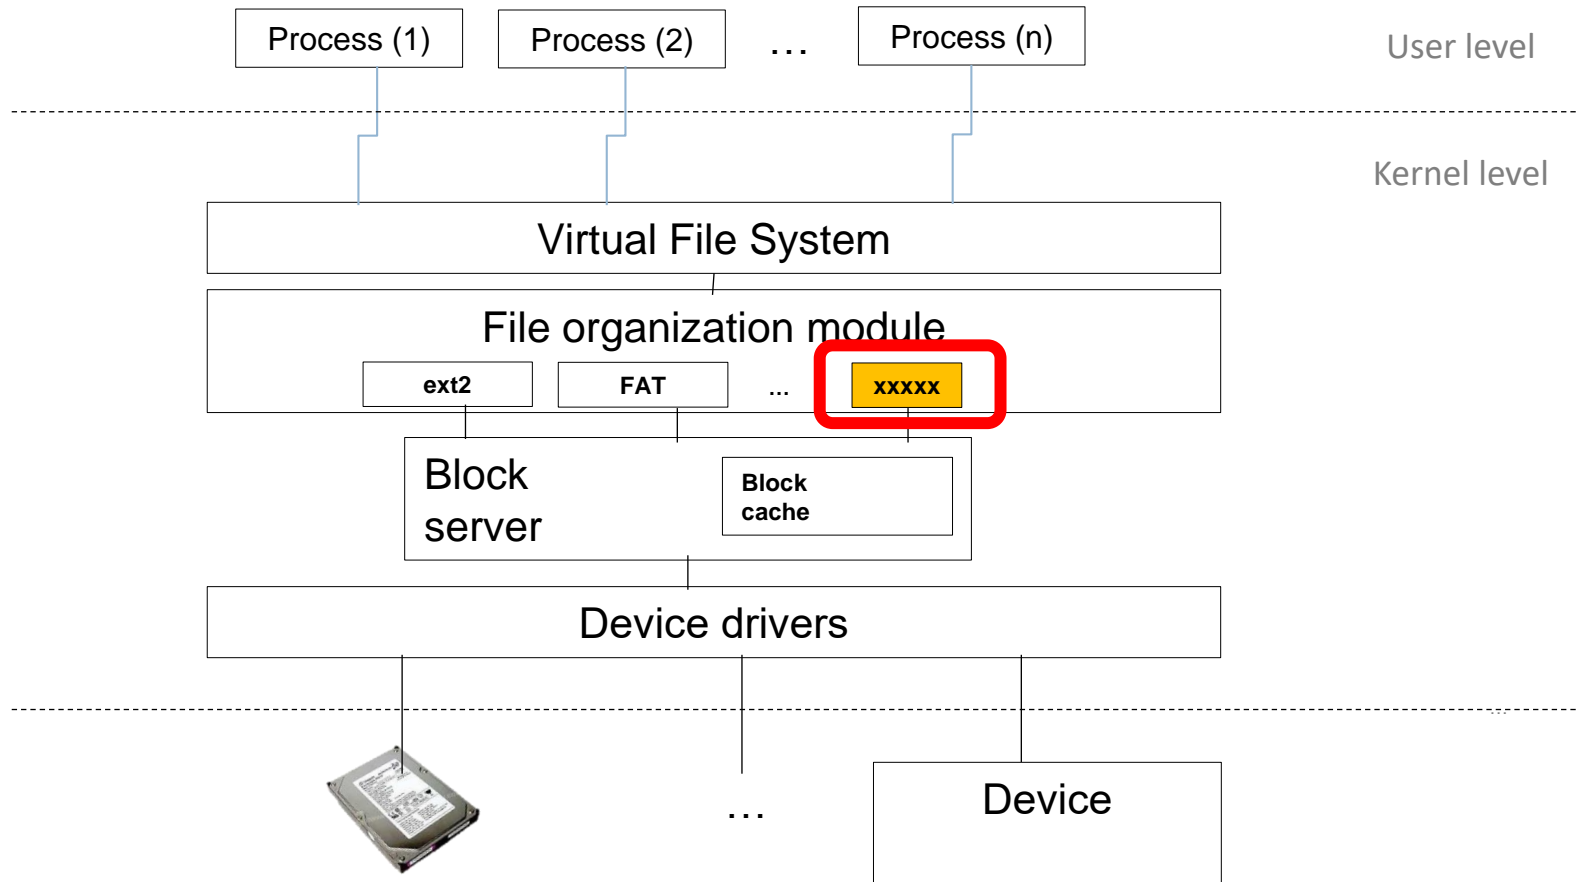
■ If a block is not in the cache, it is necessary to make room for it

■ Replacement policy

Destination (related to architecture)... file system design and implementation

35

Alejandro Calderón Mateos 

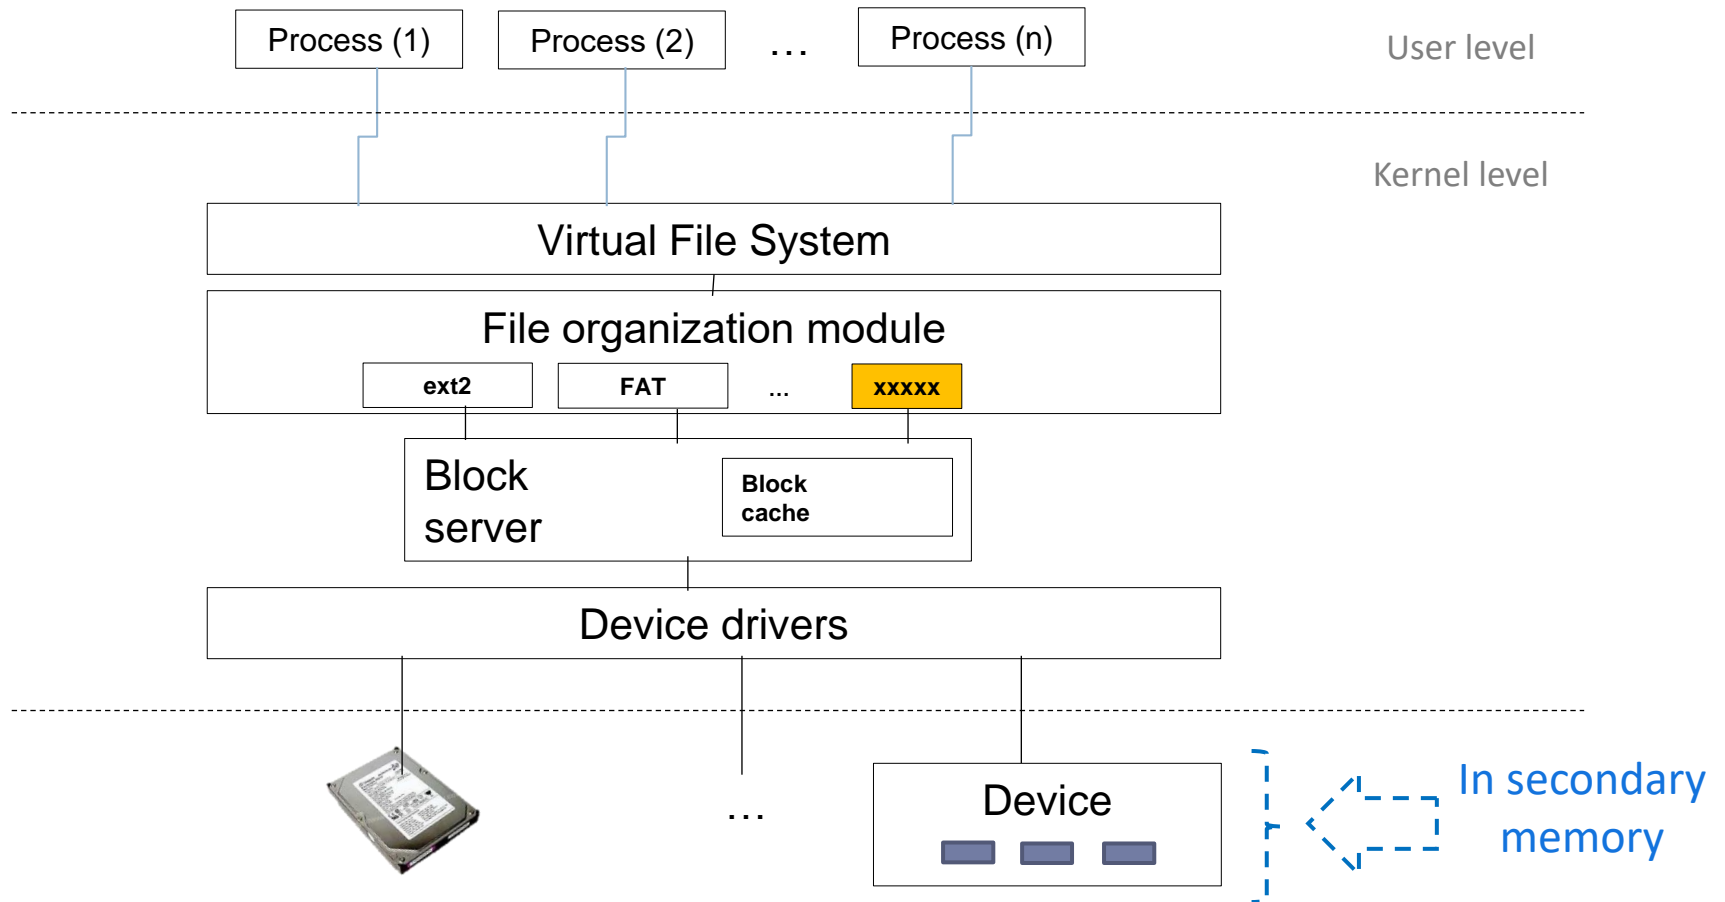


Aspects to be design (related to architecture)...

(1) Data structures on disk...

36

Alejandro Calderón Mateos 

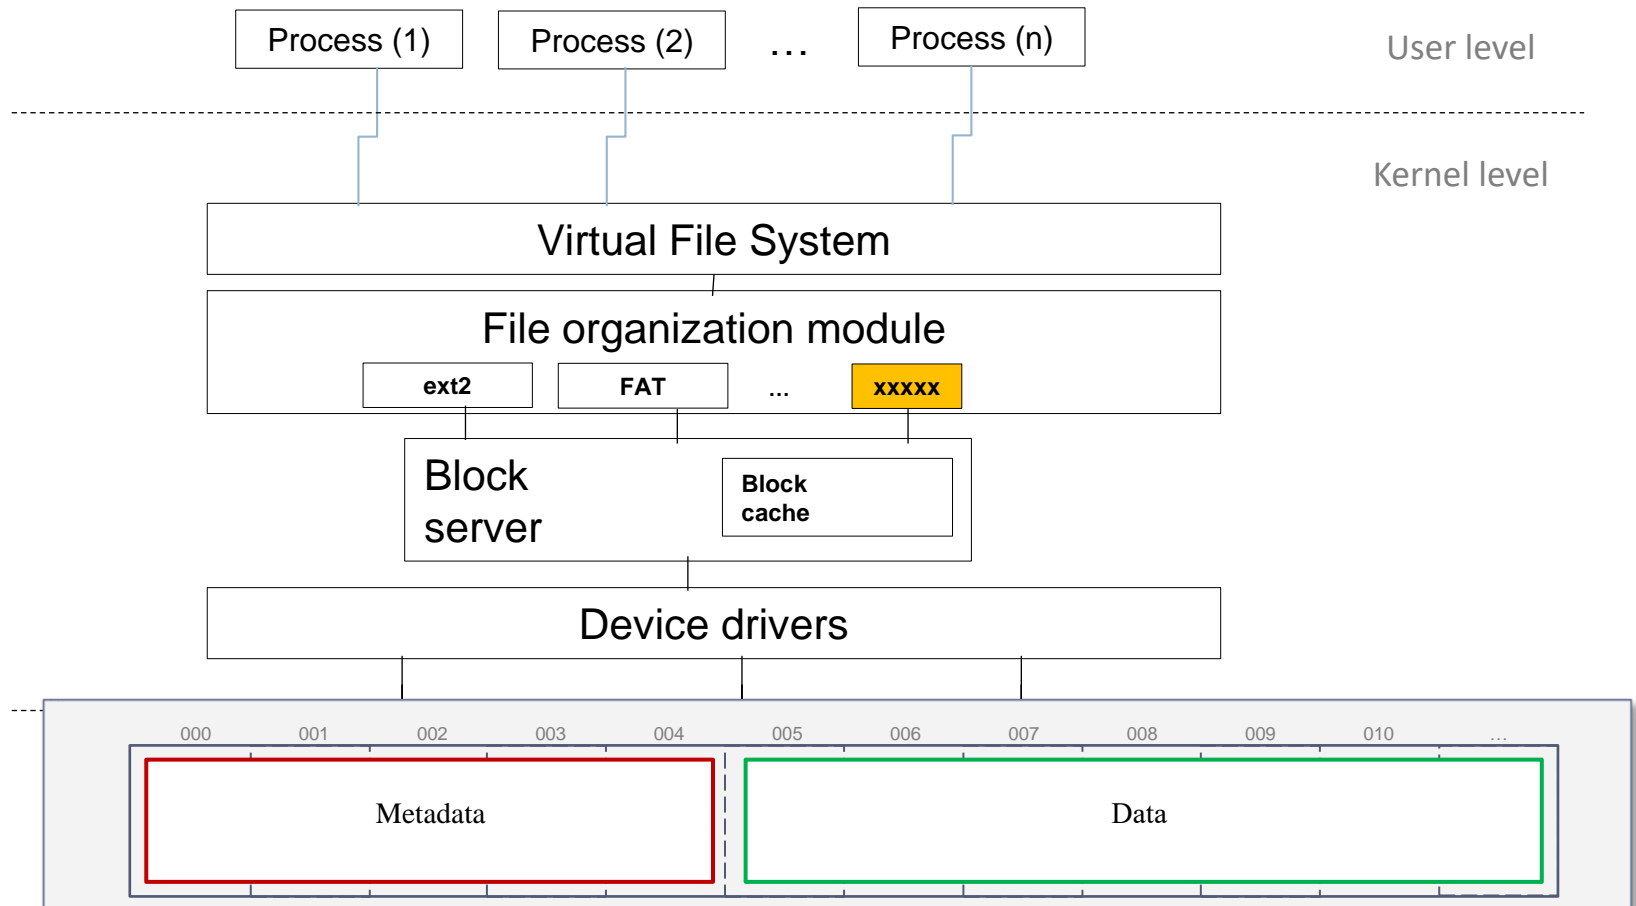


Aspects to be design (related to architecture)...

(1) Data structures on disk...

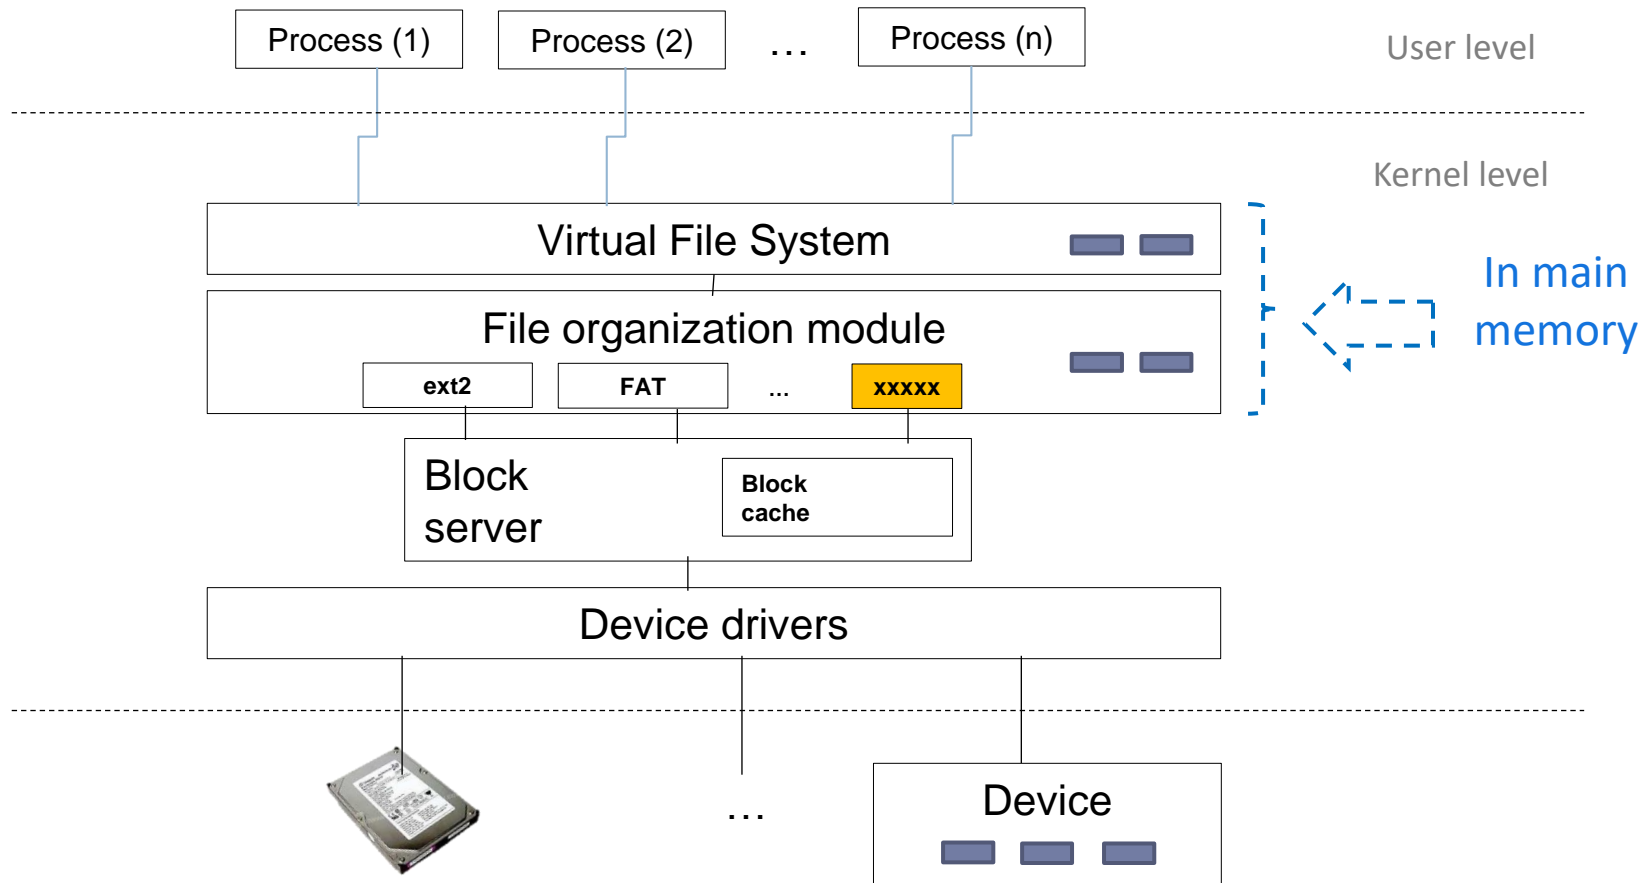
37

Alejandro Calderón Mateos



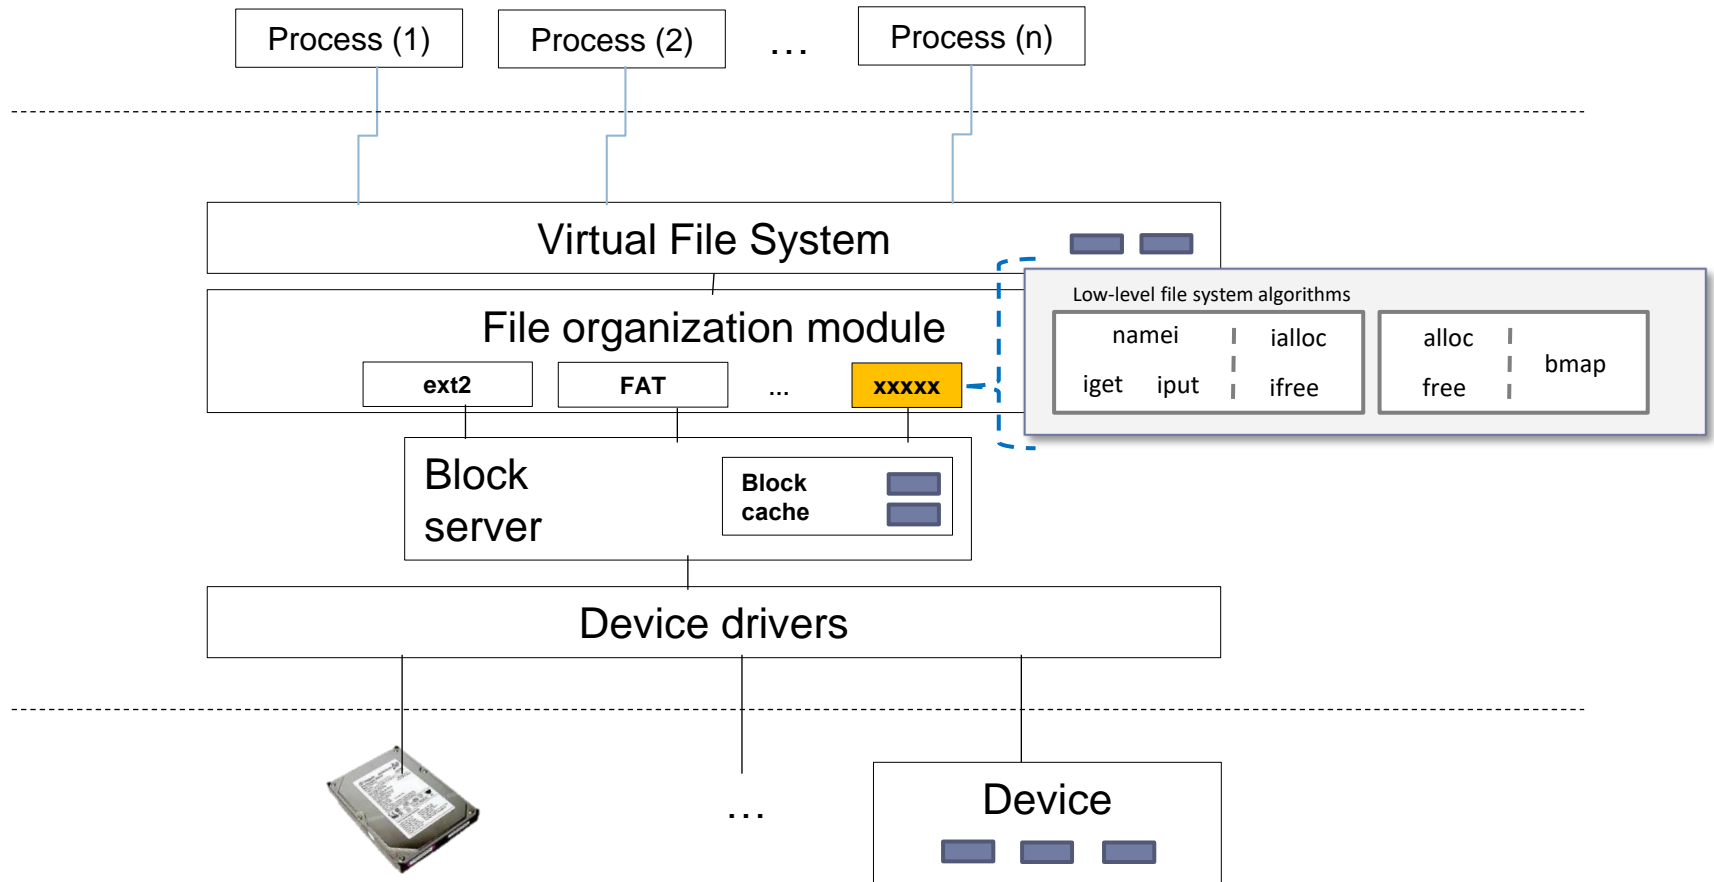
Aspects to be design (related to architecture)...

(2) Data structures in memory...



Aspects to be design (related to architecture)...

(3a) Management of disk/memory structures ...

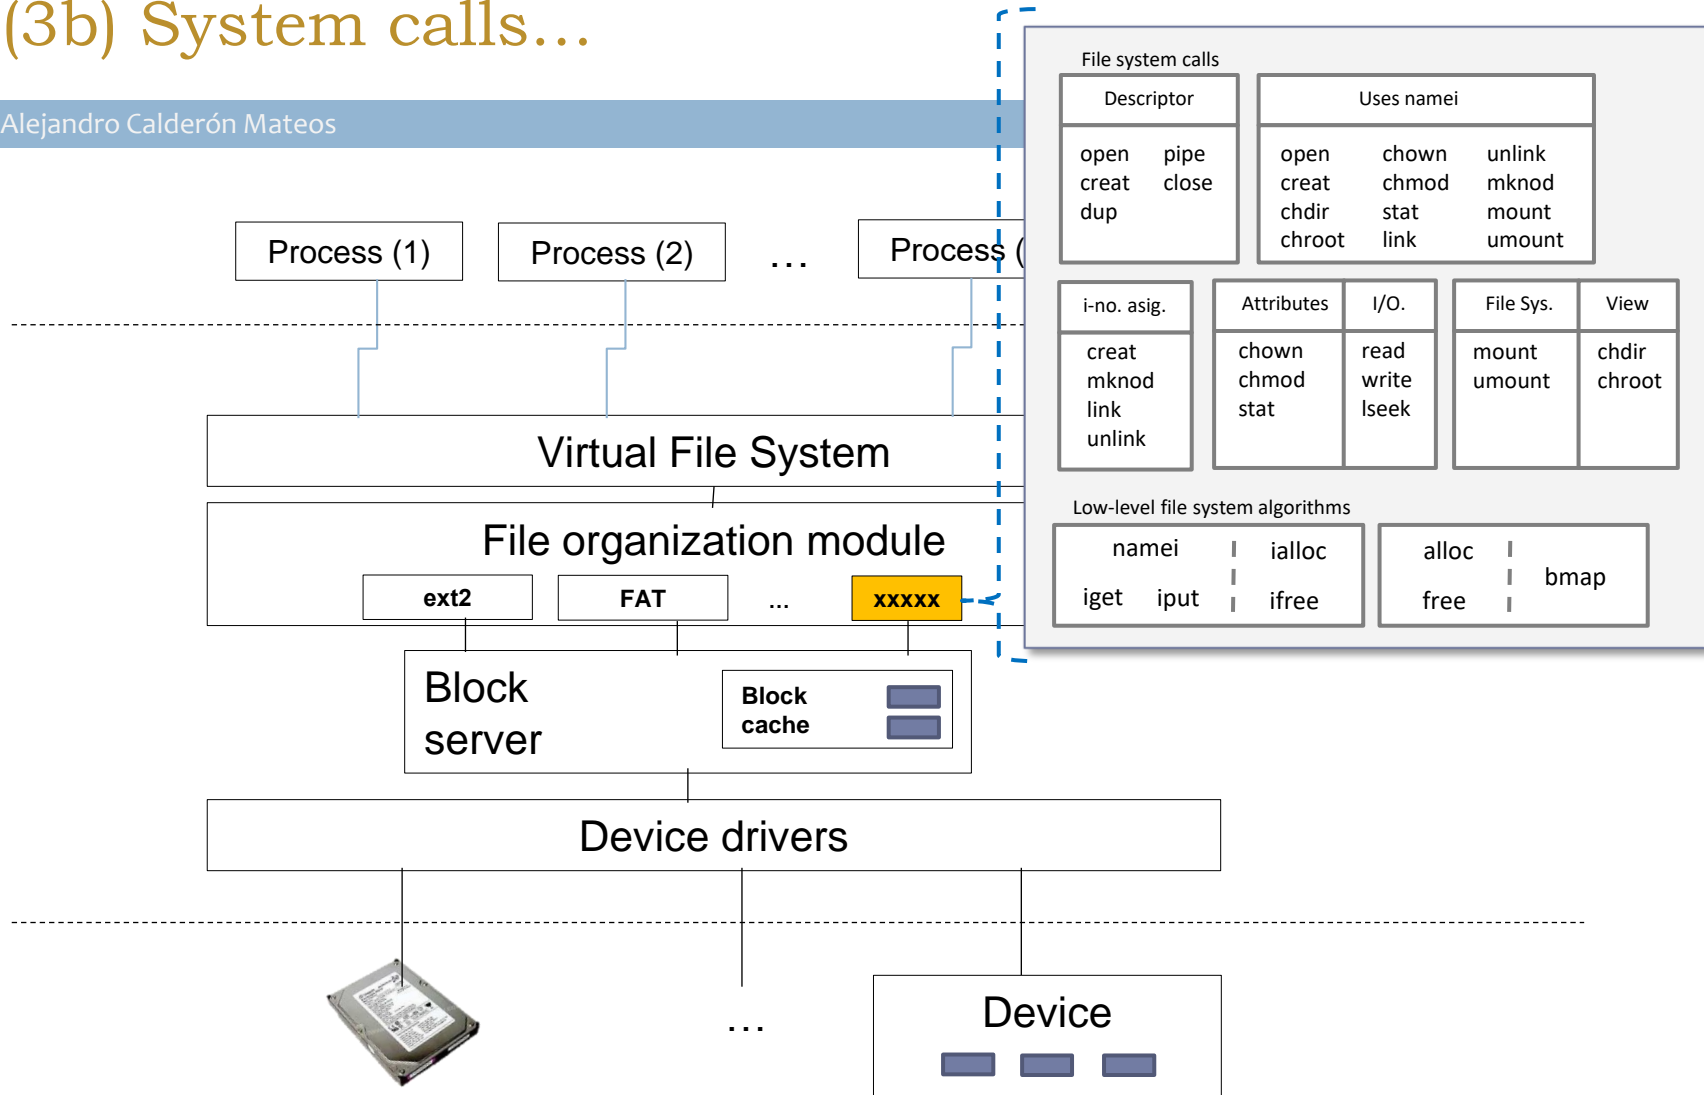


Aspects to be design (related to architecture)...

(3b) System calls...

40

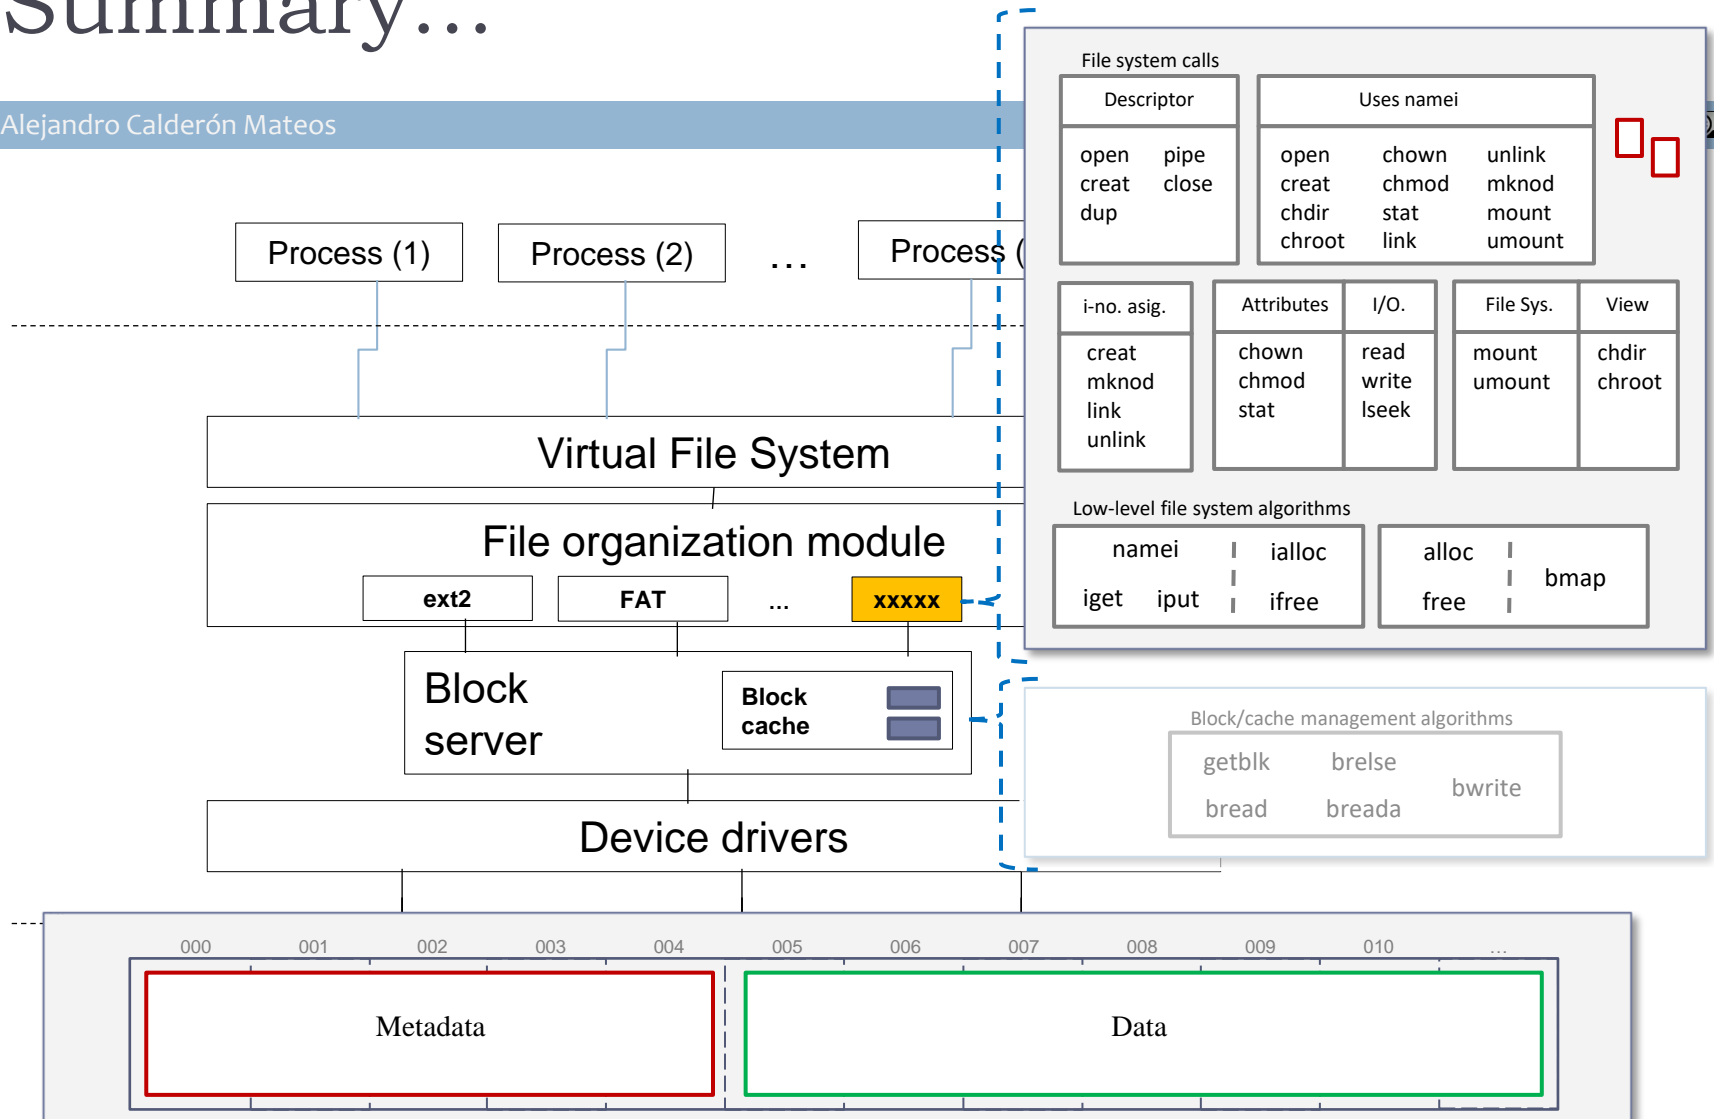
Alejandro Calderón Mateos



Summary...

41

Alejandro Calderón Mateos



Simplified summary...

42

http://www.ual.es/~acorral/DSO/Tema_4.pdf

Alejandro Calderón Mateos 

File system calls

| Descriptor | Uses namei | i-no. asig. | Attributes | I/O. | File Sys. | View |
|---------------------------------|--|----------------------------------|------------------------|------------------------|-----------------|-----------------|
| open pipe creat close dup | open chown unlink creat chmod mknod chdir stat mount chroot link umount | creat mknod link unlink | chown chmod stat | read write lseek | mount umount | chdir chroot |

Low-level file system algorithms

| | | | |
|-----------|--------|-------|------|
| namei | ialloc | alloc | |
| iget iput | ifree | free | bmap |

d-entries

mounted

file r/w pointers

open files

i-nodes in use

file system modules

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|

| | | | | | | | | | | | |
|------------|-------------|---------------------|----------------|--------------------------|-----|-----|-----|-----|-----|-----|-----|
| 000 | 001 | 002 | 003 | 004 | 005 | 006 | 007 | 008 | 009 | 010 | ... |
| Boot block | Super-block | Resource allocation | 000 i-nodes | 001 002 003 004 | | | | | | | |

ARCOS @ UC3M

Sistemas Operativos – Files, directorios y systems de ficheros

Elements to be analyzed (1, 2, 3a y 3b)

43

http://www.ual.es/~acorral/DSO/Tema_4.pdf

Alejandro Calderón Mateos 

File system calls

| Descriptor | Uses namei | i-no. asig. | Attributes | I/O. | File Sys. | View |
|---------------------------------|--|----------------------------------|------------------------|------------------------|-----------------|-----------------|
| open pipe creat close dup | open chown unlink creat chmod mknod chdir stat mount chroot link umount | creat mknod link unlink | chown chmod stat | read write lseek | mount umount | chdir chroot |

Low-level file system algorithms

| | | | |
|-----------|--------|-------|------|
| namei | ialloc | alloc | bmap |
| iget iput | ifree | free | |

file r/w pointers

d-entries

open files

mounted

i-nodes in use

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|

file system modules

| | | | | | | | | | | | |
|------------|-------------|---------------------|--|-----|-----|-----|-----|-----|-----|-----|-----|
| 000 | 001 | 002 | 003 | 004 | 005 | 006 | 007 | 008 | 009 | 010 | ... |
| Boot block | Super-block | Resource allocation | 000 001 002 003 004 i-nodes | | | | | | | | |

(1) Data structures on disk...

44

http://www.ual.es/~acorra/DSO/Tema_4.pdf

Alejandro Calderón Mateos 

File system calls

| Descriptor | Uses namei | i-no. asig. | Attributes | I/O. | File Sys. | View |
|---------------------------------|--|----------------------------------|------------------------|------------------------|-----------------|-----------------|
| open pipe creat close dup | open chown unlink creat chmod mknod chdir stat mount chroot link umount | creat mknod link unlink | chown chmod stat | read write lseek | mount umount | chdir chroot |

Low-level file system algorithms

| | | | |
|-----------|--------|-------|------|
| namei | ialloc | alloc | |
| iget iput | ifree | free | bmap |

d-entries

mounted

file r/w pointers


open files

i-nodes in use

file system modules

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|



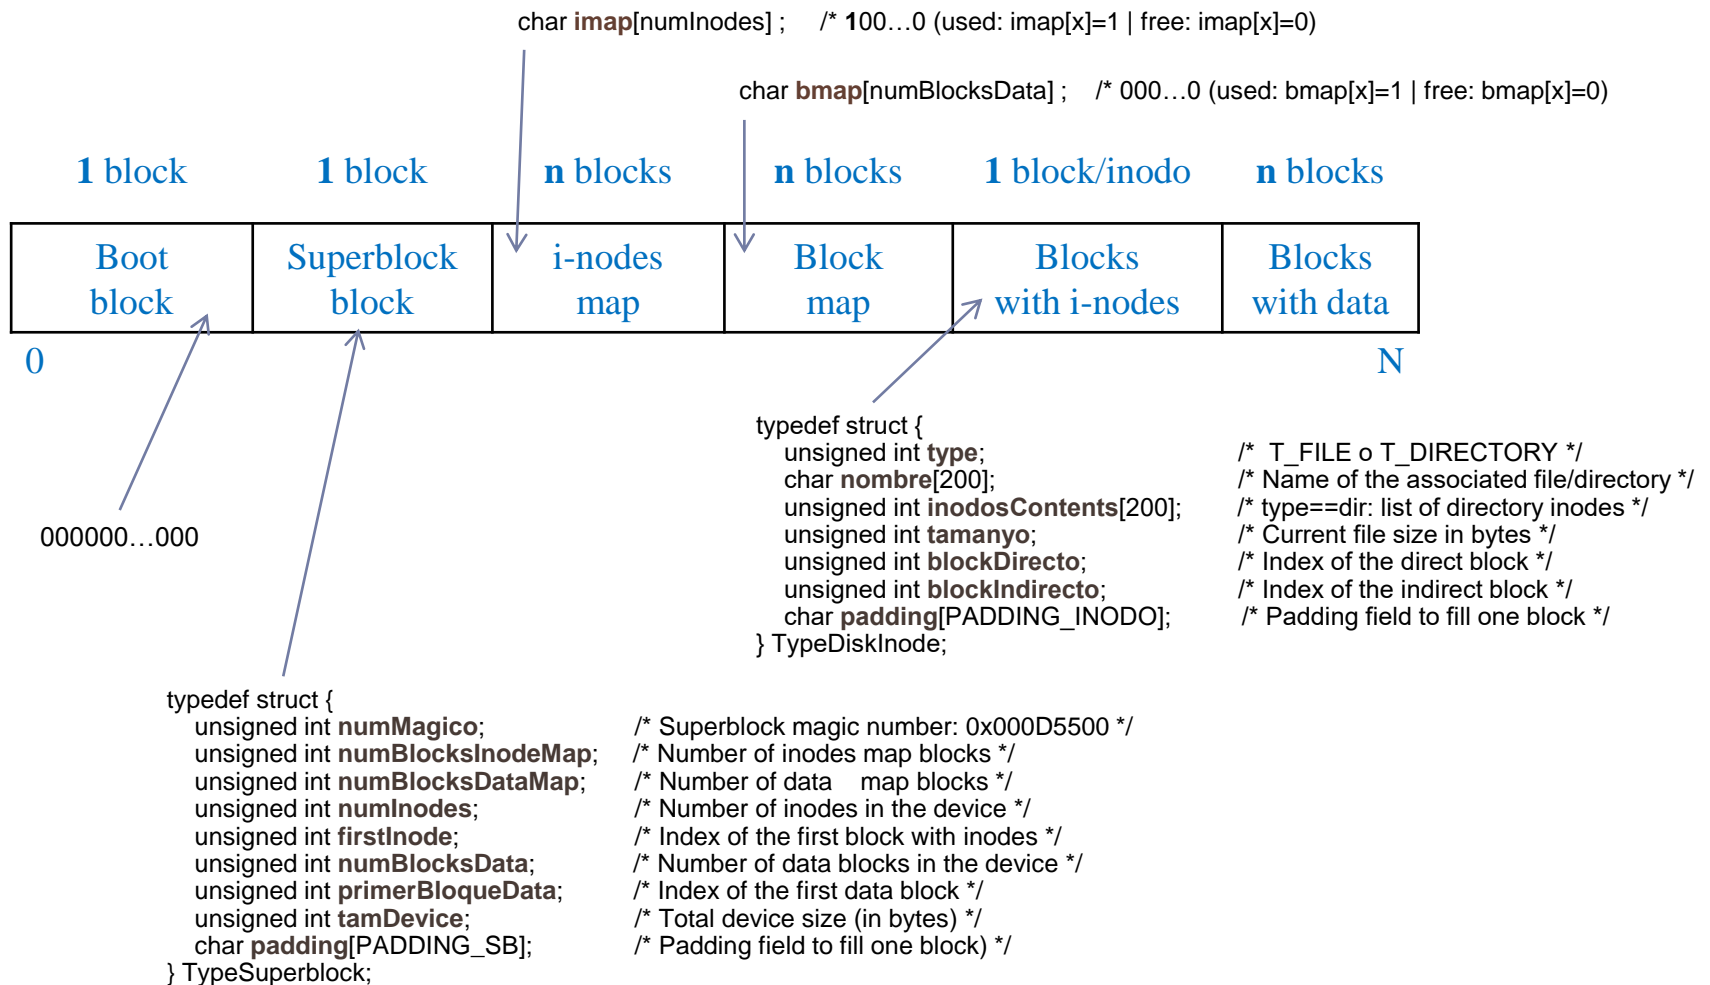
Example of disk organization

<https://github.com/acaldero/nanofs>



45

Alejandro Calderón Mateos



ARCOS @ UC3M

Sistemas Operativos – Files, directorios y systems de ficheros

(2) Data structures on memory...

46

http://www.ual.es/~acorral/DSO/Tema_4.pdf

Alejandro Calderón Mateos 

File system calls


| Descriptor | Uses namei | i-no. asig. | Attributes | I/O. | File Sys. | View |
|---------------------------------|--|----------------------------------|------------------------|------------------------|-----------------|-----------------|
| open pipe creat close dup | open chown unlink creat chmod mknod chdir stat mount chroot link umount | creat mknod link unlink | chown chmod stat | read write lseek | mount umount | chdir chroot |

Low-level file system algorithms

| | | | |
|-----------|--------|-------|------|
| namei | ialloc | alloc | |
| iget iput | ifree | free | bmap |

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|

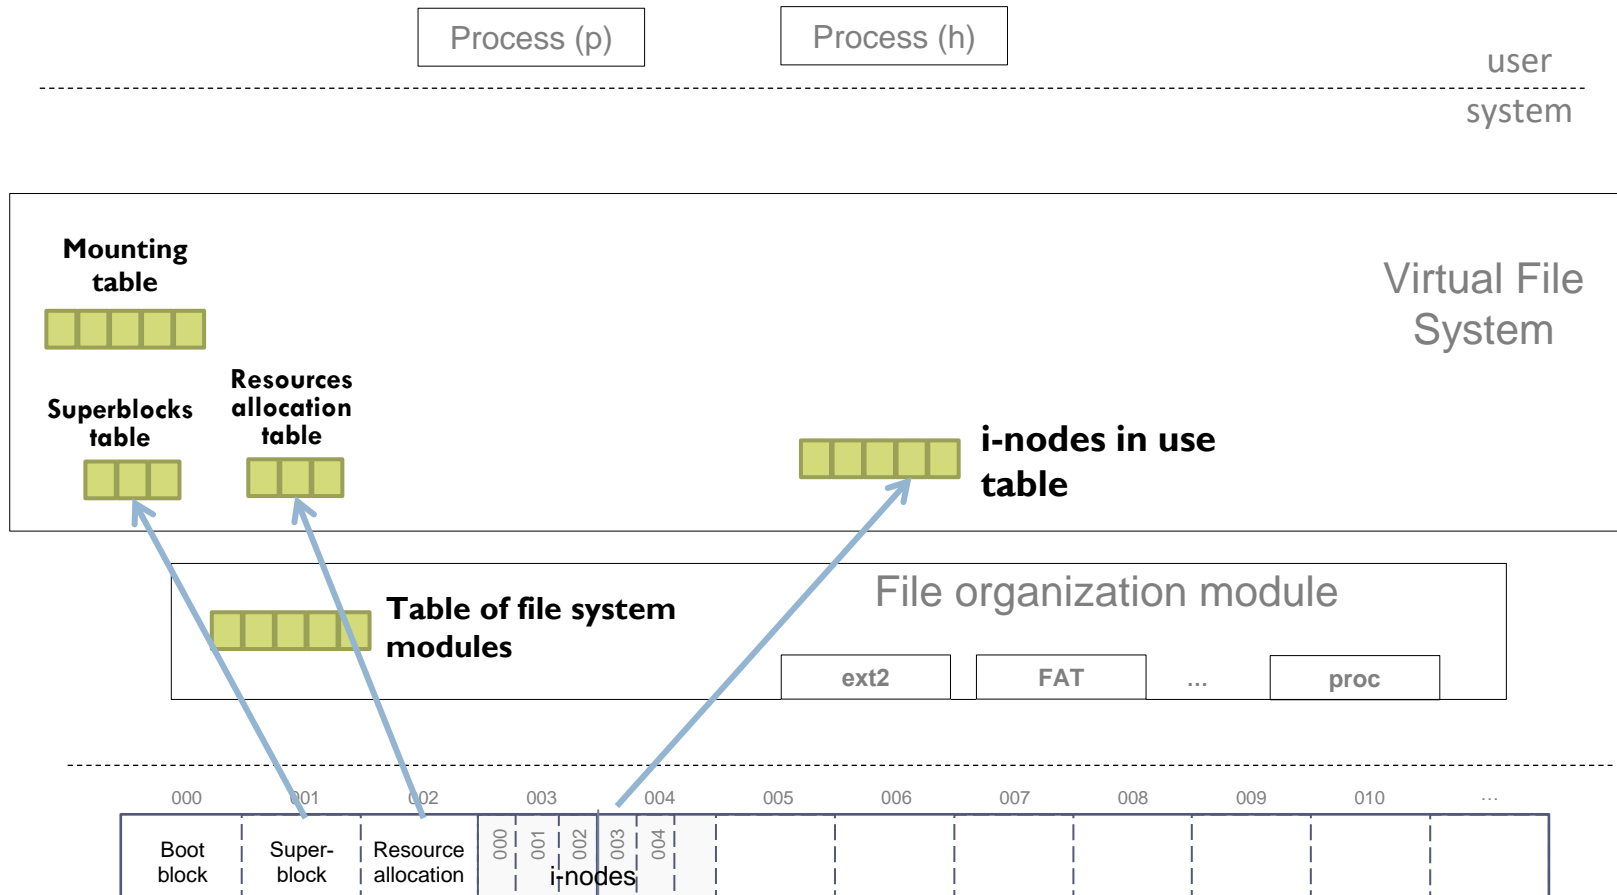


Main management structures

main metadata on disk...

47

Alejandro Calderón Mateos 

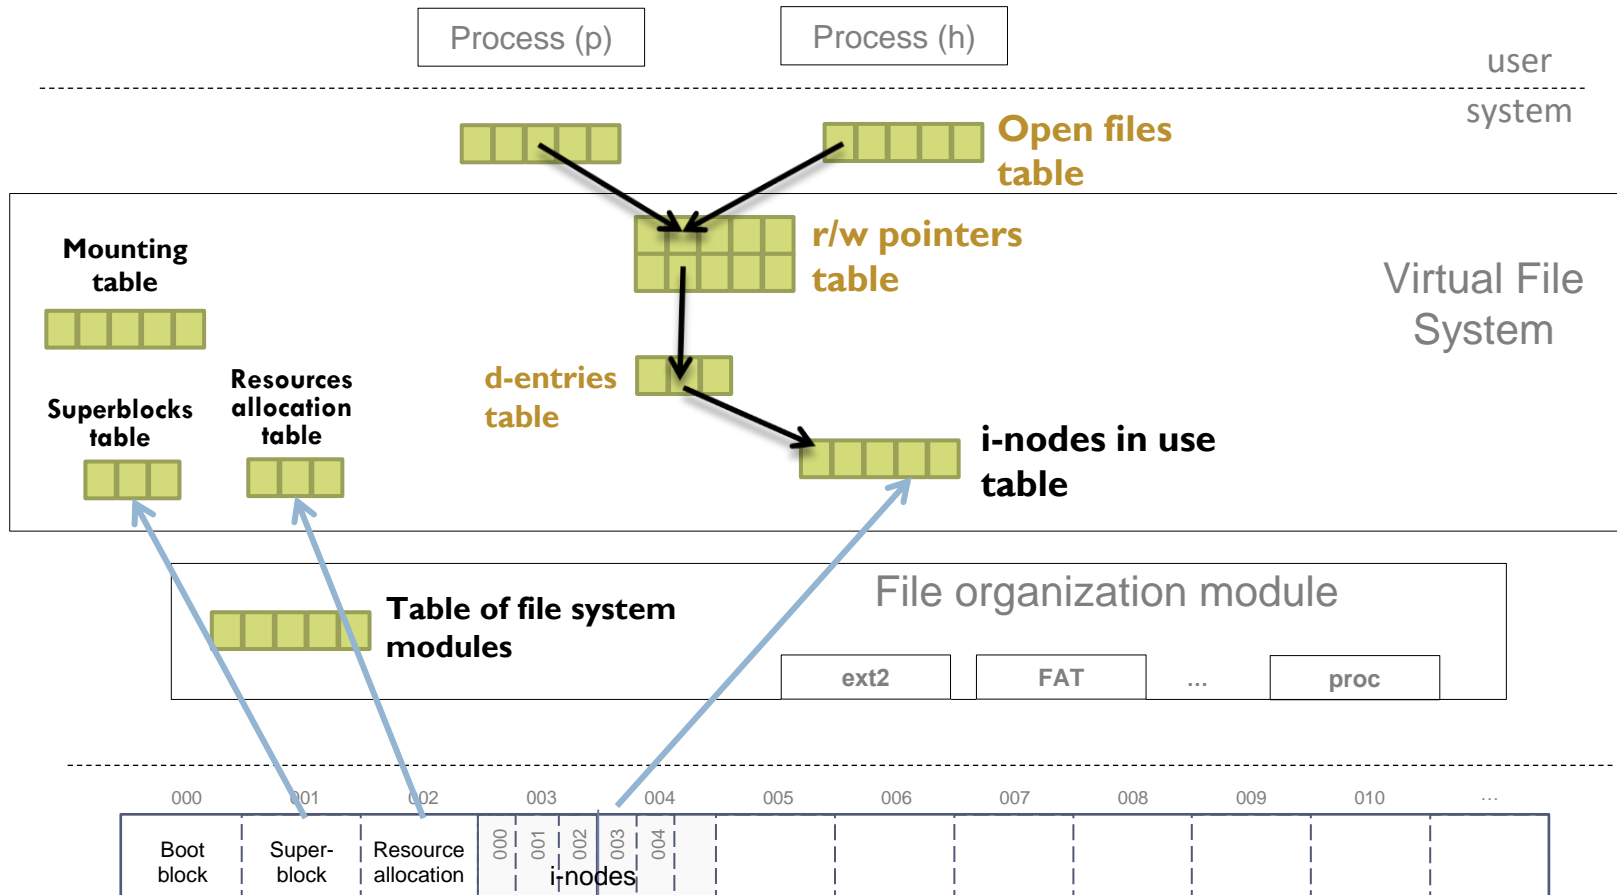


Main management structures

main metadata on disk... + 3

48

Alejandro Calderón Mateos

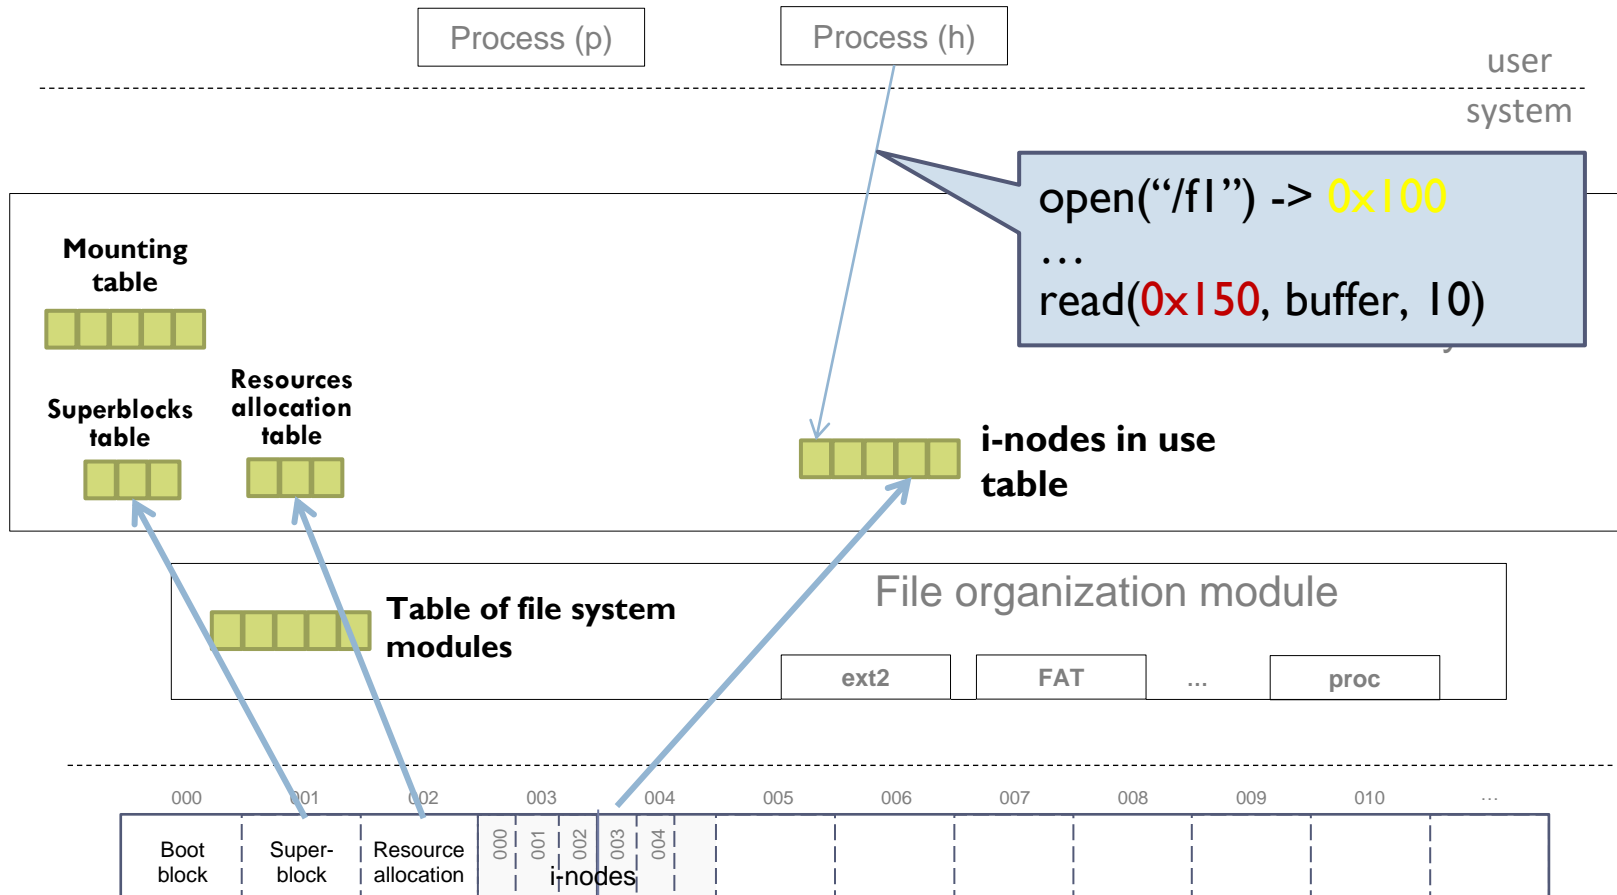


Main management structures

secure API interface?

49

Alejandro Calderón Mateos

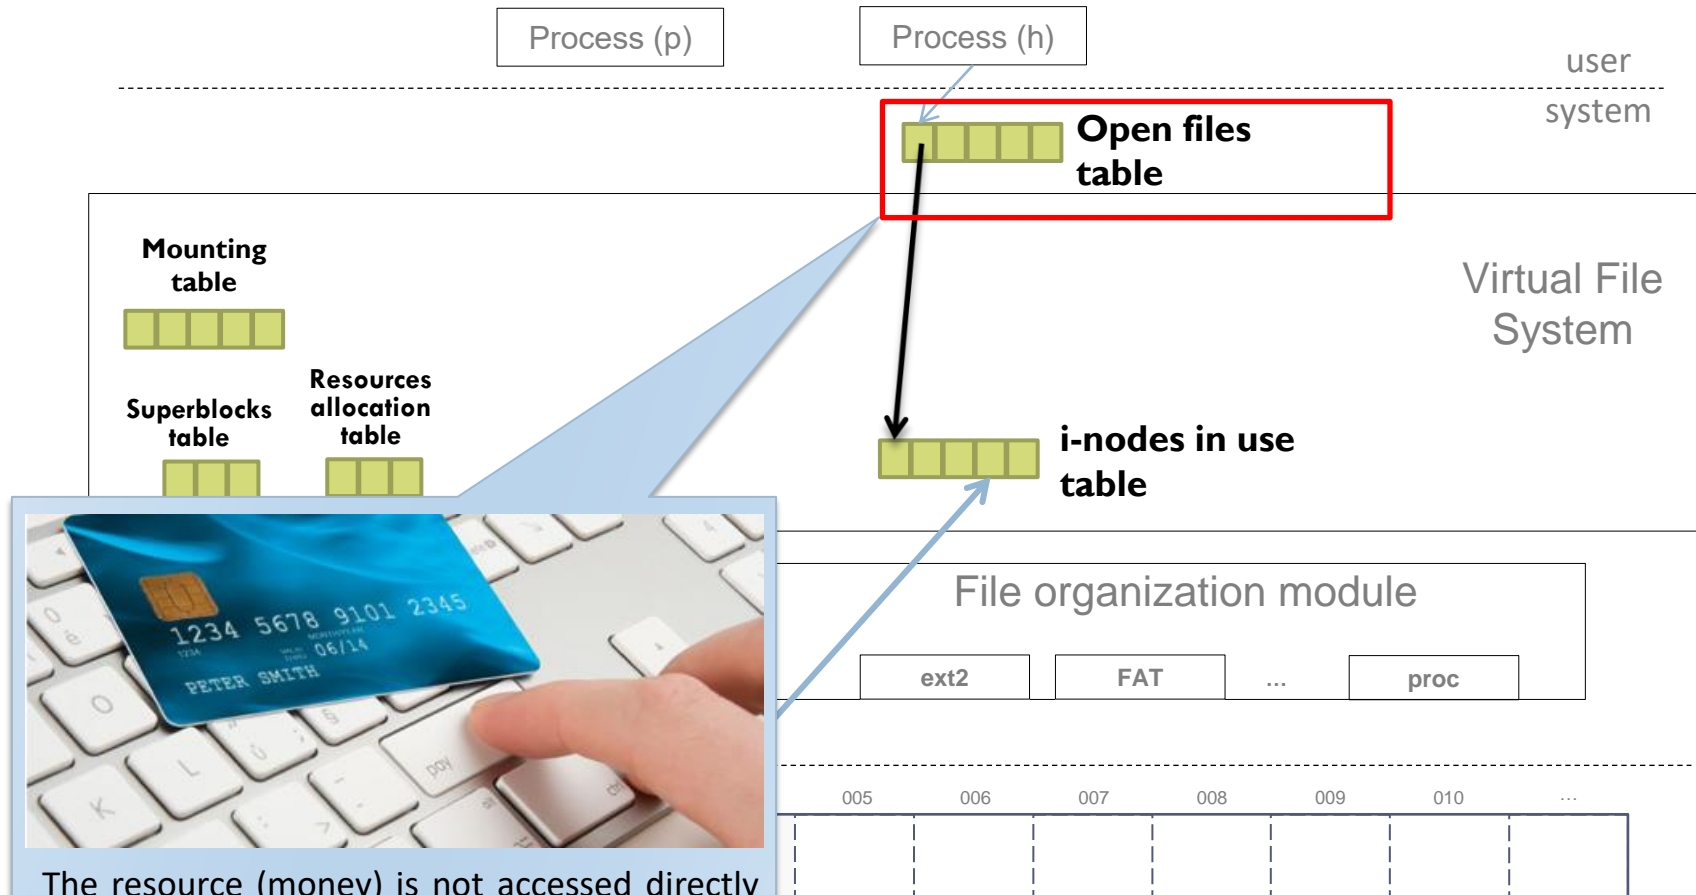


Main management structures

open files table: secure interface

50

Alejandro Calderón Mateos 



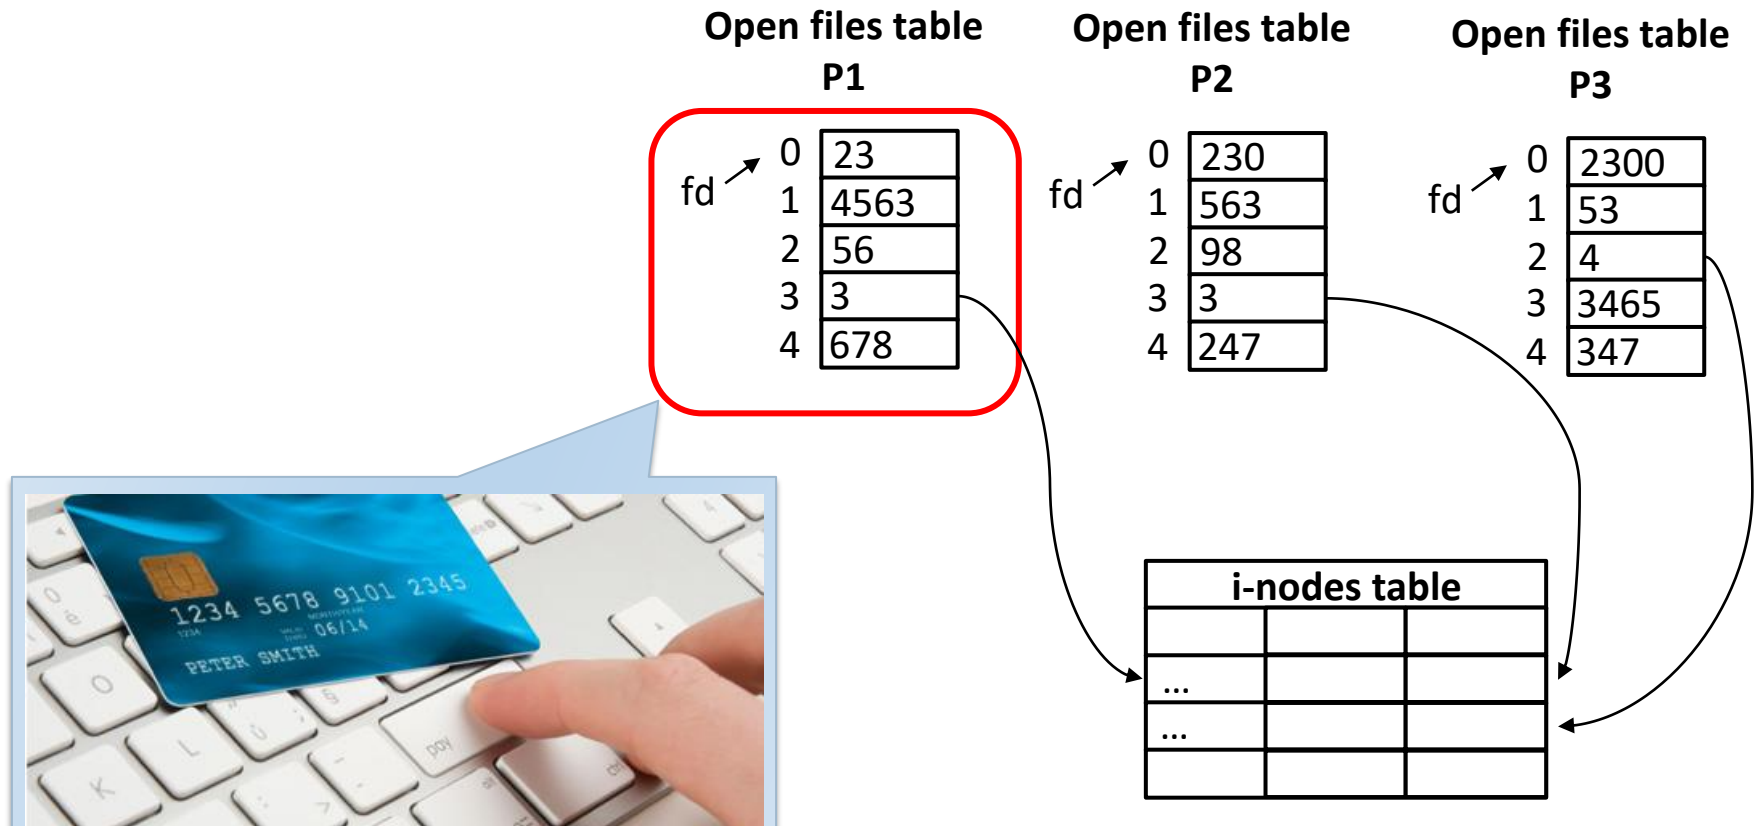
The resource (money) is not accessed directly but through a descriptor (# card).

Main management structures

open files table: secure interface

51

Alejandro Calderón Mateos



The resource (money) is not accessed directly but through a descriptor (# card).

Main management structures

open files table: secure interface

52

Alejandro Calderón Mateos



Open files table
P1

| | |
|--------|------|
| fd → 0 | 23 |
| 1 | 4563 |
| 2 | 56 |
| 3 | 3 |
| 4 | 678 |

Open files table
P2

| | |
|--------|-----|
| fd → 0 | 230 |
| 1 | 563 |
| 2 | 98 |
| 3 | 3 |
| 4 | 247 |

Open files table
P3

| | |
|--------|------|
| fd → 0 | 2300 |
| 1 | 53 |
| 2 | 4 |
| 3 | 3465 |
| 4 | 347 |

| i-nodes table | | |
|---------------|--|--|
| | | |
| ... | | |
| ... | | |
| | | |

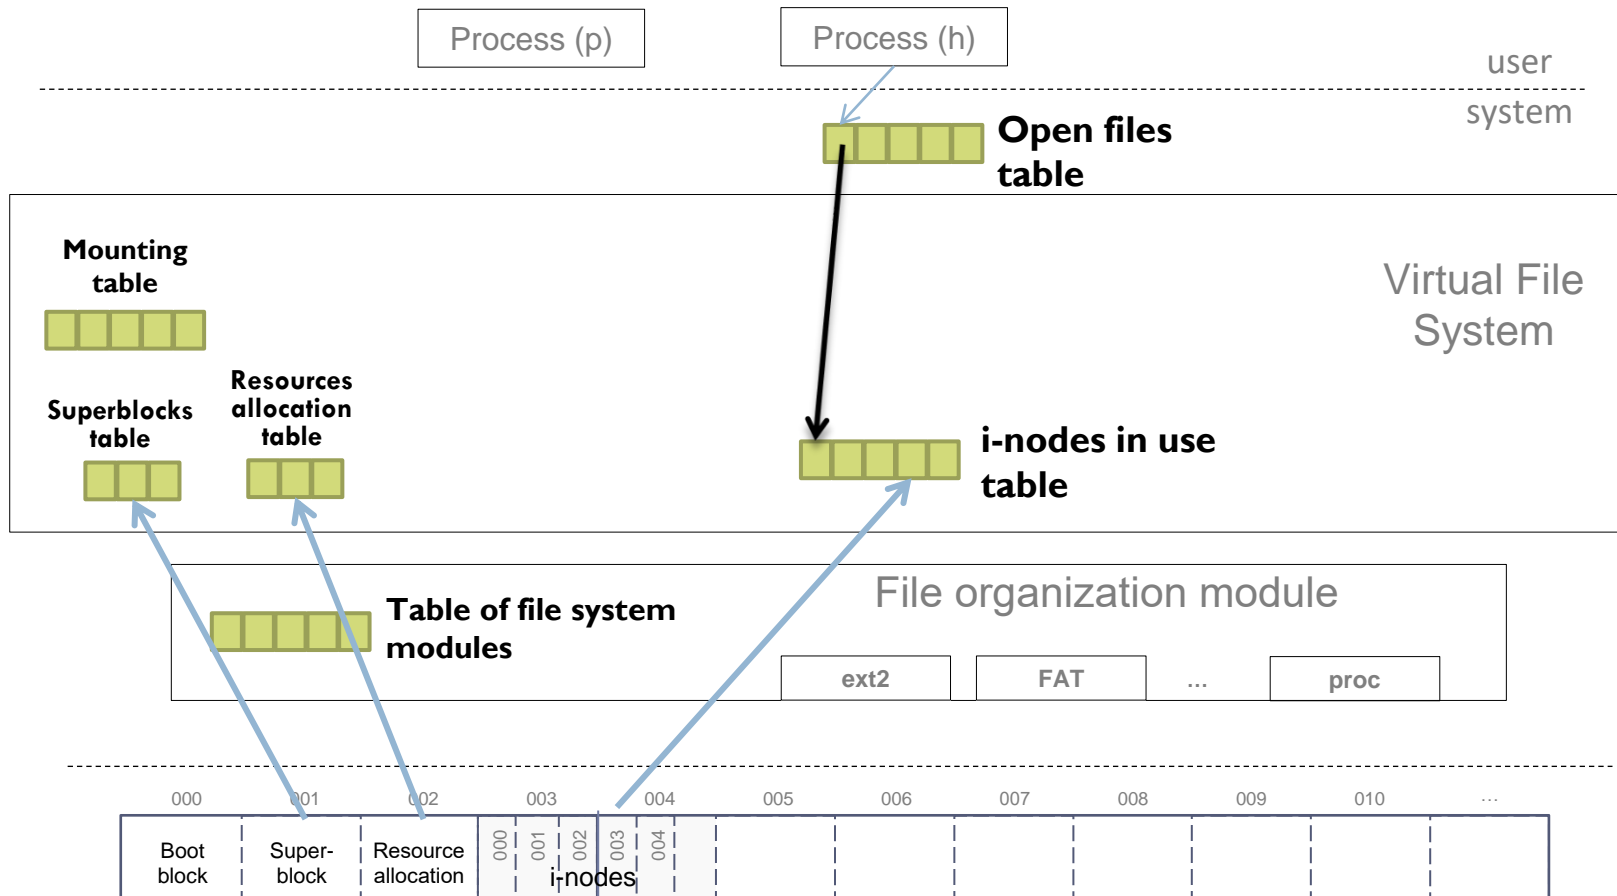
- Table **included in the BCP of the process.**
 - When `fork()` is performed, it is duplicated.
- Table with **one entry per open file.**
 - 0, 1 and 2 used by default.
- **Number of rows limits the maximum number of open files per process.**
- The table is **filled in orderly fashion:**
 - `open/creat/dup`: search first free entry.
 - `close`: marks entry as free.

Main management structures

open files table: secure interface

53

Alejandro Calderón Mateos

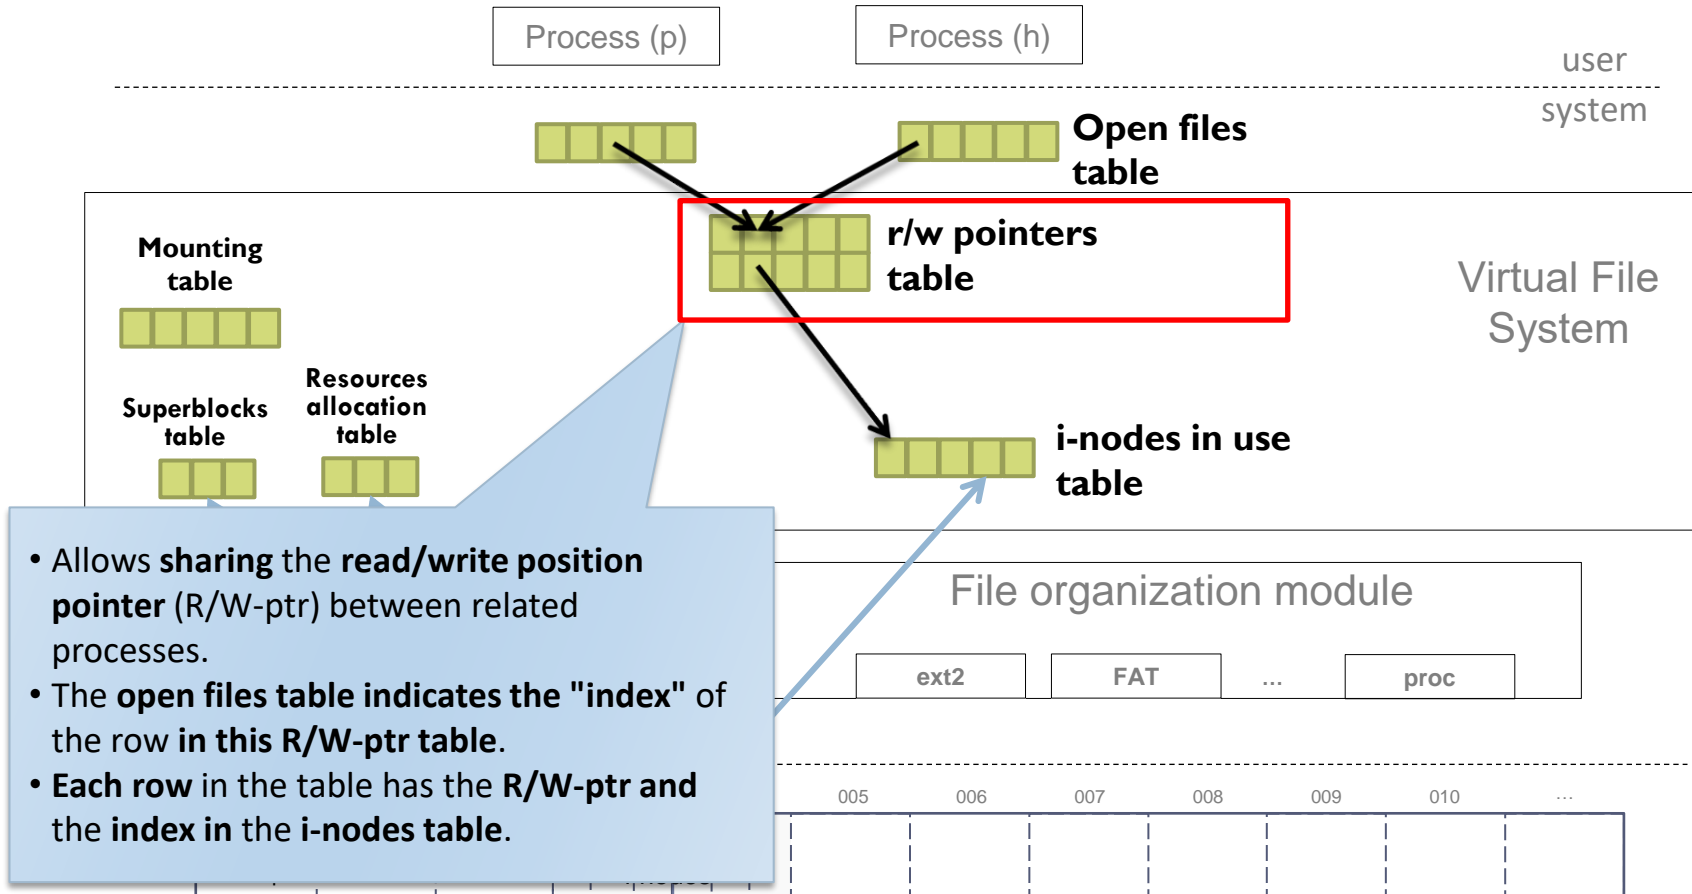


Main management structures

table of file r/w pointers: sharing r/w ptr.

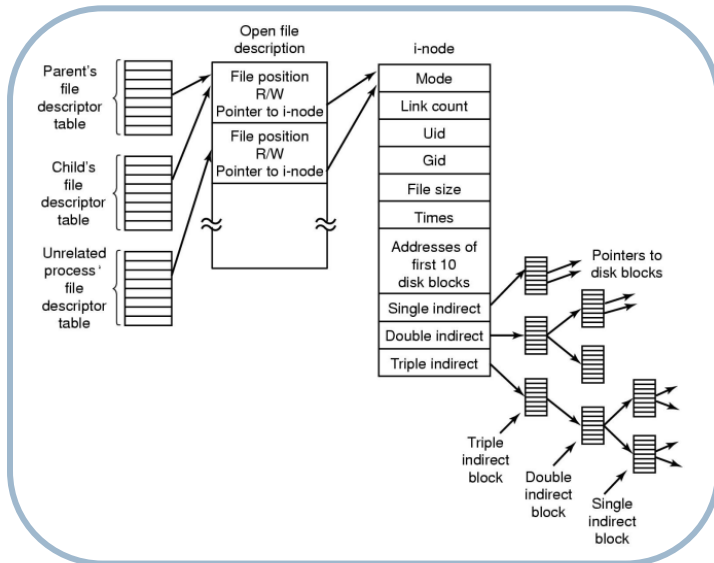
54

Alejandro Calderón Mateos



Main management structures

table of file r/w pointers: sharing r/w ptr.



Open files table

P1

| | |
|--------|------|
| fd → 0 | 23 |
| 1 | 4563 |
| 2 | 56 |
| 3 | 3 |
| 4 | 678 |

Open files table

P2

| | |
|--------|-----|
| fd → 0 | 230 |
| 1 | 563 |
| 2 | 98 |
| 3 | 3 |
| 4 | 247 |

Open files table

P3

| | |
|--------|------|
| fd → 0 | 2300 |
| 1 | 53 |
| 2 | 4 |
| 3 | 3465 |
| 4 | 347 |

I-nodes table

| | | |
|--|--|--|
| | | |
| | | |
| | | |
| | | |
| | | |

| i-Node | Offset |
|--------|--------|
| 92 | 345 |
| 92 | 5678 |
| | |

Intermediate table of i-nodes and offsets

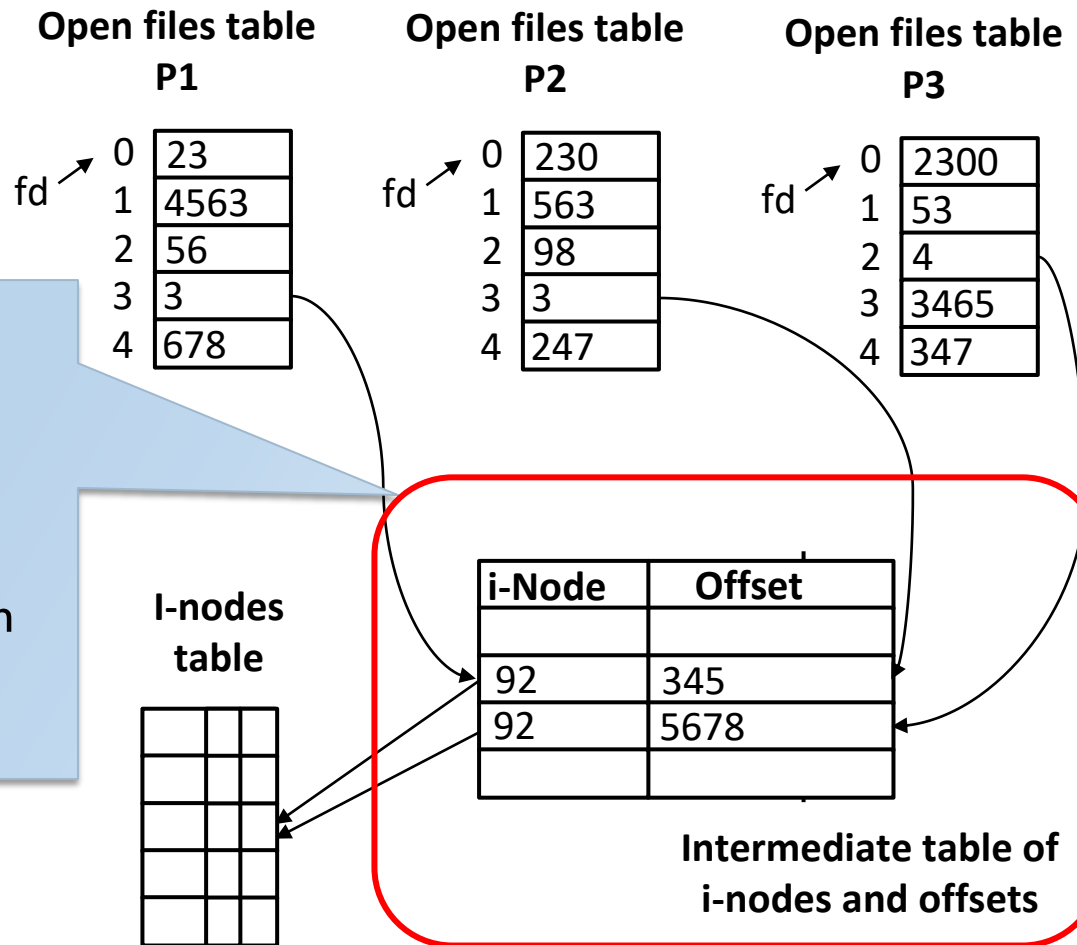
Main management structures

table of file r/w pointers: sharing r/w ptr.

56

Sistemas operativos: una visión aplicada (© J. Carrete et al.)

Alejandro Calderón Mateos



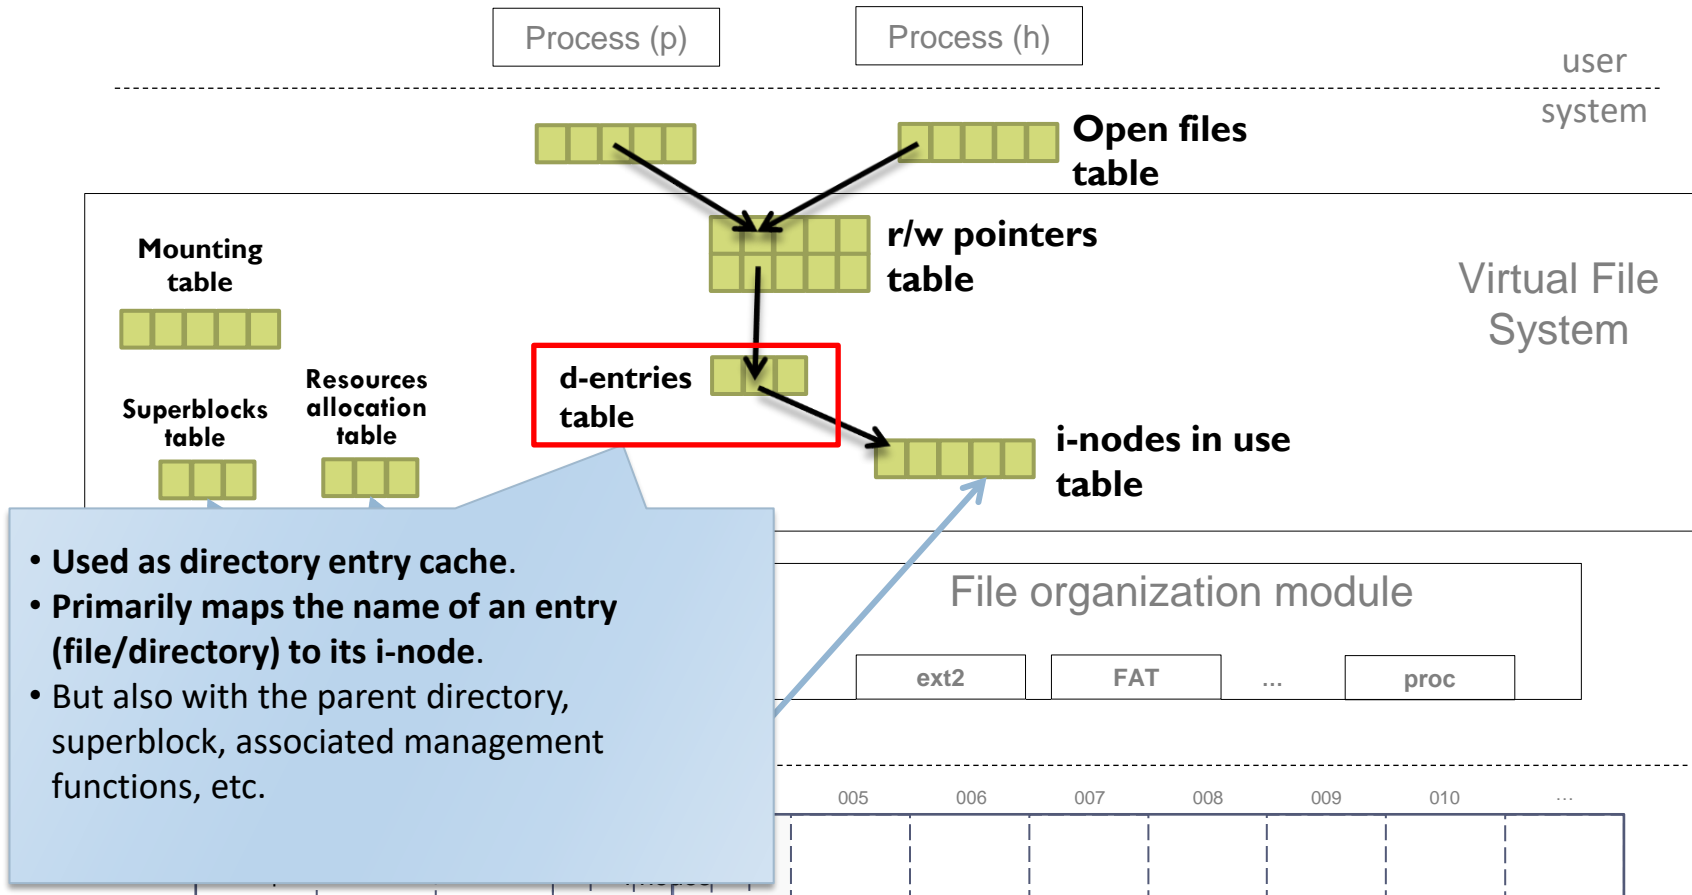
- FILP table (FILE Pointer)
- Between the descriptor table and (usually) the i-node table.
- Saves (mainly) the file position pointer.

Main management structures

d-entries table: working with directories

57

Alejandro Calderón Mateos 

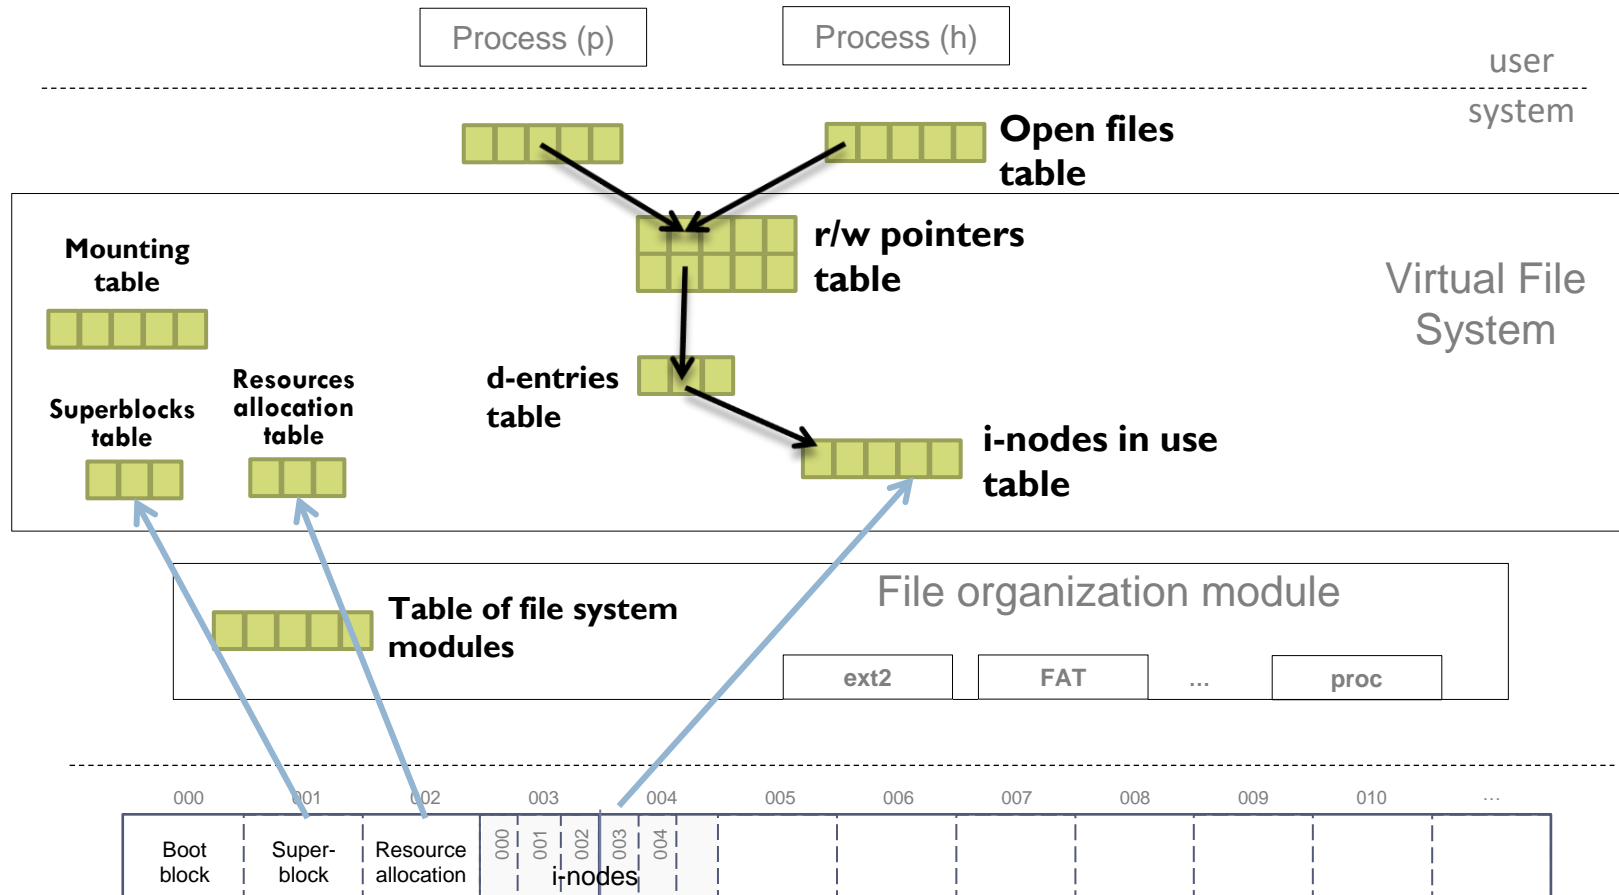


Main management structures

summary of the main data structures in memory

58

Alejandro Calderón Mateos 



Example of memory organization...

<https://github.com/acaldero/nanofs>



```
// Information read from the disk
TypeSuperblock sblocks [1] ;
char imap [numInodo] ;
char bmap [numBlocksData] ;
TypeDiskInode inodos [numInodo] ;

// Extra support information
struct {
    int posicion;
    int abierto;
} inodos_x [numInodo] ;

...
```

(3a) Management of disk/memory structures ...

60

http://www.ual.es/~acorral/DSO/Tema_4.pdf

Alejandro Calderón Mateos 

File system calls

| Descriptor | Uses namei | i-no. asig. | Attributes | I/O. | File Sys. | View |
|---------------------------------|--|----------------------------------|------------------------|------------------------|-----------------|-----------------|
| open pipe creat close dup | open chown unlink creat chmod mknod chdir stat mount chroot link umount | creat mknod link unlink | chown chmod stat | read write lseek | mount umount | chdir chroot |

Low-level file system algorithms

| | | | |
|-----------|--------|-------|------|
| namei | ialloc | alloc | |
| iget iput | ifree | free | bmap |

d-entries

mounted

file r/w pointers

open files

i-nodes in use

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|

file system modules



ARCOS @ UC3M

Sistemas Operativos – Files, directorios y systems de ficheros

Example of management routines

i-nodes

61

<http://www.buet.ac.bd/iict/iictcourses/ict6005/lecture9.ppt>

Alejandro Calderón Mateos



- ▶ **namei**: converts a path to the associated i-node.
- ▶ **iget**: returns an i-node from the i-node table and if not present, reads it from secondary memory, adds it to the i-node table and returns it.
- ▶ **iput**: releases an i-node from the i-node table, and if necessary, updates it in secondary memory.
- ▶ **ialloc**: allocates an i-node to a file.
- ▶ **ifree**: releases an i-node previously assigned to a file.

Low-level file system algorithms

| | | | |
|-------|--------|-------|------|
| namei | ialloc | alloc | |
| iget | iput | free | bmap |

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|

d-entries



mounted



file r/w pointers

open files

i-nodes in use



file system modules

| | | | | | | | | | | | |
|------------|-------------|---------------------|---------|-----|-----|-----|-----|-----|-----|-----|-----|
| 000 | 001 | 002 | 003 | 004 | 005 | 006 | 007 | 008 | 009 | 010 | ... |
| Boot block | Super-block | Resource allocation | 000 | 001 | 002 | 003 | 004 | | | | |
| | | | i-nodes | | | | | | | | |

Example of management routines

blocks

62

<http://www.buet.ac.bd/iict/iictcourses/ict6005/lecture9.ppt>

Alejandro Calderón Mateos



- ▶ **bmap**: calculates the disk block associated with a file offset. Translates logical addresses (file offset) to physical addresses (disk block).
- ▶ **alloc**: allocates a block to a file.
- ▶ **free**: releases a block previously assigned to a file.

Low-level file system algorithms

| | | | |
|-------|--------|-------|------|
| namei | ialloc | alloc | |
| iget | iput | free | bmap |
| | ifree | | |

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|

d-entries

mounted

file r/w pointers

open files

i-nodes in use

file system modules

| | | | | | | | | | | | |
|------------|-------------|---------------------|---------|-----|-----|-----|-----|-----|-----|-----|-----|
| 000 | 001 | 002 | 003 | 004 | 005 | 006 | 007 | 008 | 009 | 010 | ... |
| Boot block | Super-block | Resource allocation | 000 | 001 | 002 | 003 | 004 | | | | |
| | | | i-nodes | | | | | | | | |

Example: ialloc and alloc

<https://github.com/acaldero/nanofs>



63

<http://lsi.ugr.es/~jlgarrid/so2/pdf/tem2-1-2.pdf>

Alejandro Calderón Mateos 

```
int ialloc ( void )
{
    // buscar un i-nodo libre
    for (int=0; i<sblocks[0].numInodes; i++)
    {
        if (imap[i] == 0) {
            // inodo ocupado ahora
            imap[i] = 1;
            // valores por defecto en el i-nodo
            memset(&(inodos[i]),0,
                sizeof(TypeDiskInode));
            // devolver identificador de i-nodo
            return i;
        }
    }

    return -1;
}
```

```
int alloc ( void )
{
    char b[BLOCK_SIZE];

    for (int=0; i<sblocks[0].numBlocksData; i++)
    {
        if (bmap[i] == 0) {
            // block ocupado ahora
            bmap[i] = 1;
            // valores por defecto en el block
            memset(b, 0, BLOCK_SIZE);
            bwrite(DISK, sblocks[0].primerBloqueData + i, b);
            // devolver identificador del block
            return i;
        }
    }

    return -1;
}
```

Example: ifree and free

<https://github.com/acaldero/nanofs>



64

http://mycsvtunotes.weebly.com/uploads/1/0/1/7/10174835/unix_unit4.pdf

Alejandro Calderón Mateos 

```
int ifree ( int inodo_id )
{
    // comprobar validez de inodo_id
    if (inodo_id > sblocks[0].numInodes)
        return -1;

    // liberar i-nodo
    imap[inodo_id] = 0;

    return -1;
}
```

```
int free ( int block_id )
{
    // comprobar validez de block_id
    if (block_id > sblocks[0].numBlocksData)
        return -1;

    // liberar block
    bmap[block_id] = 0;

    return -1;
}
```


Example: namei and bmap

<https://github.com/acaldero/nanofs>



65

http://mycsvtunotes.weebly.com/uploads/1/0/1/7/10174835/unix_unit4.pdf

Alejandro Calderón Mateos 

```
int namei ( char *fname )
{
    // buscar i-nodo con nombre <fname>
    for (int=0; i<sblocks[0].numInodes; i++)
    {
        if (! strcmp(inodos[i].nombre, fname))
            return i;
    }

    return -1;
}
```

```
int bmap ( int inodo_id, int offset )
{
    int b[BLOCK_SIZE/4];

    // comprobar validez de inodo_id
    if (inodo_id > sblocks[0].numInodes)
        return -1;

    // block de datos asociado
    if (offset < BLOCK_SIZE)
        return inodos[inodo_id].blockDirecto;
    if (offset < BLOCK_SIZE*BLOCK_SIZE/4) {
        bread(DISK, sblocks[0].primerBloqueData +
            inodos[inodo_id].blockIndirecto, b);
        offset = (offset - BLOCK_SIZE) / BLOCK_SIZE;
        return b[offset] ;
    }

    return -1;
}
```

(3b) System calls...

66

http://www.ual.es/~acorra/DSO/Tema_4.pdf


Alejandro Calderón Mateos 

File system calls

| Descriptor | Uses namei | i-no. asig. | Attributes | I/O. | File Sys. | View |
|---------------------------------|--|----------------------------------|------------------------|------------------------|-----------------|-----------------|
| open pipe creat close dup | open chown unlink creat chmod mknod chdir stat mount chroot link umount | creat mknod link unlink | chown chmod stat | read write lseek | mount umount | chdir chroot |

Low-level file system algorithms

| | | | |
|-----------|--------|-------|------|
| namei | ialloc | alloc | bmap |
| iget iput | ifree | free | |



Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|



Example

sys. calls

- ▶ **open**: locates the i-node associated with the path of the file, ...
- ▶ **read**: locates the data block, read data block, ...
- ▶ **write**: locate the data block, write data block, ...
- ▶ ...

67

http://mycsvtunotes.weebly.com/uploads/1/0/1/7/10174835/unix_unit4.pdf

Juanjo Calderón Mateos

File system calls

| Descriptor | Uses namei | i-no. asig. | Attributes | I/O. | File Sys. | View |
|---------------------------------|--|----------------------------------|------------------------|------------------------|-----------------|-----------------|
| open pipe creat close dup | open chown unlink creat chmod mknod chdir stat mount chroot link umount | creat mknod link unlink | chown chmod stat | read write lseek | mount umount | chdir chroot |

Low-level file system algorithms

| | | | |
|-----------|--------|-------|------|
| namei | ialloc | alloc | |
| iget iput | ifree | free | bmap |

d-entries

mounted

file r/w pointers

open files

i-nodes in use

file system modules

Block/cache management algorithms

| | | | | |
|--------|--------|-------|--------|--------|
| getblk | brelse | bread | breada | bwrite |
|--------|--------|-------|--------|--------|



ARCOS @ UC3M

Sistemas Operativos – Files, directorios y systems de ficheros

Example: mount

<https://github.com/acaldero/nanofs>



```
int mount ( void )
{
    // leer block 0 de disco en sblocks[0]
    bread(DISK, 1, &(sblocks[0]) );

    // leer los blocks para el mapa de i-nodes
    for (int=0; i<sblocks[0].numBlocksInodeMap; i++)
        bread(DISK, 2+i, ((char *)imap + i*BLOCK_SIZE) );

    // leer los blocks para el mapa de blocks de datos
    for (int=0; i<sblocks[0].numBlocksDataMap; i++)
        bread(DISK, 2+i+sblocks[0].numBlocksInodeMap, ((char *)bmap + i*BLOCK_SIZE);

    // leer los i-nodes a memoria
    for (int=0; i<(sblocks[0].numInodes*sizeof(TypeDiskInode)/BLOCK_SIZE); i++)
        bread(DISK, i+sblocks[0].firstInode, ((char *)inodos + i*BLOCK_SIZE);

    return 1;
}
```

Example: umount

<https://github.com/acaldero/nanofs>



```
int umount ( void )
{
    // escribir block 0 de sblocks[0] a disco
    bwrite(DISK, 1, &(sblocks[0]) );

    // escribir los blocks para el mapa de i-nodes
    for (int=0; i<sblocks[0].numBlocksInodeMap; i++)
        bwrite(DISK, 2+i, ((char *)imap + i*BLOCK_SIZE) );

    // escribir los blocks para el mapa de blocks de datos
    for (int=0; i<sblocks[0].numBlocksDataMap; i++)
        bwrite(DISK, 2+i+sblocks[0].numBlocksInodeMap, ((char *)bmap + i*BLOCK_SIZE);

    // escribir los i-nodes a disco
    for (int=0; i<(sblocks[0].numInodes*sizeof(TypeDiskInode)/BLOCK_SIZE); i++)
        bwrite(DISK, i+sblocks[0].firstInode, ((char *)inodos + i*BLOCK_SIZE);

    return 1;
}
```

Example: mkfs

<https://github.com/acaldero/nanofs>



```
int mkfs ( void )
{
    // inicializar a los valores por defecto del superblock, mapas e i-nodes
    sblocks[0].numMagico = 1234; // ayuda a comprobar que se haya creado por nuestro mkfs
    sblocks[0].numInodes = numInodo;
    ...
    for (int=0; i<sblocks[0].numInodes; i++)
        imap[i] = 0; // free
    for (int=0; i<sblocks[0].numBlocksData; i++)
        bmap[i] = 0; // free
    for (int=0; i<sblocks[0].numInodes; i++)
        memset(&(inodos[i]), 0, sizeof(TypeDiskInode) );

    // to write the default file system into disk
    umount();

    return 1;
}
```

Example: open and close

<https://github.com/acaldero/nanofs>



71

http://mycsvtunotes.weebly.com/uploads/1/0/1/7/10174835/unix_unit4.pdf

Alejandro Calderón Mateos 

```
int open ( char *nombre )
{
    int inodo_id ;

    inodo_id = namei(nombre) ;
    if (inodo_id < 0)
        return inodo_id ;

    inodos_x[inodo_id].posicion = 0;
    inodos_x[inodo_id].abierto  = 1;

    return inodo_id;
}
```

```
int close ( int fd )
{
    if (fd < 0)
        return fd ;

    inodos_x[fd].posicion = 0;
    inodos_x[fd].abierto  = 0;

    return 1;
}
```

Example: creat and unlink

<https://github.com/acaldero/nanofs>



72

http://mycsvtunotes.weebly.com/uploads/1/0/1/7/10174835/unix_unit4.pdf

Alejandro Calderón Mateos 

```
int creat ( char *nombre )
{
    int b_id, inodo_id ;

    inodo_id = ialloc() ;
    if (inodo_id < 0) { return inodo_id ; }
    b_id = alloc();
    if (b_id < 0) { ifree(inodo_id); return b_id ; }

    inodos[inodo_id].type = 1 ; // FICHERO
    strcpy(inodos[inodo_id].nombre, nombre);
    inodos[inodo_id].blockDirecto = b_id ;
    inodos_x[inodo_id].posicion = 0;
    inodos_x[inodo_id].abierto  = 1;

    return 1;
}
```

```
int unlink ( char * nombre )
{
    int inodo_id ;

    inodo_id = namei(nombre) ;
    if (inodo_id < 0)
        return inodo_id ;

    free(inodos[inodo_id].blockDirecto);
    memset(&(inodos[inodo_id]),
           0,
           sizeof(TypeDiskInode));
    ifree(inodo_id) ;

    return 1;
}
```


Example: read and write

<https://github.com/acaldero/nanofs>



73

http://mycsvtunotes.weebly.com/uploads/1/0/1/7/10174835/unix_unit4.pdf

Alejandro Calderón Mateos 

```
int read ( int fd, char *buffer, int size )
{
    char b[BLOCK_SIZE] ;
    int b_id ;

    if (inodos_x[fd].posicion+size > inodos[fd].size)
        size = inodos[fd].size - inodos_x[fd].posicion;
    if (size <= 0)
        return 0;

    b_id = bmap(fd, inodos_x[fd].posicion);
    bread(DISK,
        sblocks[0].primerBloqueData+b_id, b);
    memmove(buffer,
        b+inodos_x[fd].posicion, size);
    inodos_x[fd].posicion += size;

    return size;
}
```

```
int write ( int fd, char *buffer, int size )
{
    char b[BLOCK_SIZE] ;
    int b_id ;

    if (inodos_x[fd].posicion+size > BLOCK_SIZE)
        size = BLOCK_SIZE - inodos_x[fd].posicion;
    if (size <= 0)
        return 0;

    b_id = bmap(fd, inodos_x[fd].posicion);
    bread(DISK, sblocks[0].primerBloqueData+b_id, b);
    memmove(b+inodos_x[fd].posicion,
        buffer, size);
    bwrite(DISK, sblocks[0].primerBloqueData+b_id, b);
    inodos_x[fd].posicion += size;

    return size;
}
```

SISTEMAS OPERATIVOS: SISTEMAS DE FICHEROS



Files, directorios y system de ficheros